

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
# Load datasets
train_df = pd.read_excel('train_data.xlsx')
test_df = pd.read_excel('test_data.xlsx')

# Basic info
print("Train shape:", train_df.shape)
print("Test shape:", test_df.shape)
train_df.head(5)
```

Train shape: (49999, 24)
Test shape: (5931, 24)

```
Out[1]:
```

	country	device_brand	device_model	re_install	os	attribution_event_timesta
0	CZ	Blackview	Tab8	0	android	2025-01-06 15:12:35.C
1	AT	samsung	SM-S916B	0	android	2025-01-19 16:50:59.C
2	BG	HONOR	RBN-NX1	0	android	2025-01-19 08:46:42.C
3	AZ	samsung	SM-A217F	0	android	2025-01-11 09:10:43.C
4	AZ	samsung	SM-A515F	0	android	2025-01-16 16:45:16.C

5 rows x 24 columns

Output: The training set has 50,000 rows and 24 columns, and the test set has 5,931 rows (with similar columns except no churn). The first few rows of train_df confirm features like country, device_brand, os (operating system, here mostly "android"), attribution_event_timestamp (user acquisition timestamp), ecpi (acquisition cost), gameplay stats (current_gold, lvl_no, etc.), and churn (0/1).

```
In [2]: train_df.info()           # Check data types
train_df.isnull().sum()         # Count missing values per column
```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 49999 entries, 0 to 49998

Data columns (total 24 columns):

#	Column	Non-Null Count	Dtype
0	country	49999 non-null	object
1	device_brand	48745 non-null	object
2	device_model	48745 non-null	object
3	re_install	49999 non-null	int64
4	os	49999 non-null	object
5	attribution_event_timestamp	49999 non-null	object
6	ecpi	49471 non-null	float64
7	lang	49999 non-null	object
8	current_gold	49999 non-null	int64
9	totalPowerUp	49999 non-null	int64
10	bonus_cnt	49999 non-null	int64
11	duration	49999 non-null	int64
12	hint1_cnt	49999 non-null	int64
13	hint2_cnt	49999 non-null	int64
14	hint3_cnt	49999 non-null	int64
15	lvl_no	49999 non-null	int64
16	repeat_cnt	49999 non-null	int64
17	banner_impr	49999 non-null	int64
18	inter_impr	49999 non-null	int64
19	rewarded_impr	49999 non-null	int64
20	user_id	49999 non-null	int64
21	campaignid	49999 non-null	int64
22	partnerid	49999 non-null	int64
23	churn	49999 non-null	int64

dtypes: float64(1), int64(17), object(6)

memory usage: 9.2+ MB

```
Out[2]: country          0
device_brand        1254
device_model        1254
re_install          0
os                  0
attribution_event_timestamp  0
ecpi                528
lang                0
current_gold        0
totalPowerUp        0
bonus_cnt           0
duration            0
hint1_cnt           0
hint2_cnt           0
hint3_cnt           0
lvl_no              0
repeat_cnt          0
banner_impr         0
inter_impr          0
rewarded_impr       0
user_id             0
campaignid          0
partnerid           0
churn               0
dtype: int64
```

Output: Most columns are numeric or object (for categorical text fields).

We find a few columns with missing values:

-Device brand/model: ~1,254 entries (~2.5%) have missing device_brand and device_model (likely cases where device info wasn't captured).

-ECPI: 528 entries (~1.1%) have missing ecpi (acquisition cost). (It will be solved on later cells)

-Country: (In the test set, 1 missing country is observed; none missing in train).

Categorical Feature Distribution

Let's explore the categorical variables in the dataset to understand their distribution:

Country: There are 120 unique countries in the training data. The most frequent countries are, for example, France (FR), Italy (IT), Romania (RO), etc. We list the top 5 countries by number of users:

```
In [3]: print(train_df['country'].value_counts(normalize=True).head(20).mul(100).rou
```

	Percentage	Count	Cumulative
country			
FR	15.88	7938	15.88
IT	9.97	4985	25.85
RO	7.94	3969	33.79
SI	7.30	3650	41.09
BR	6.09	3045	47.18
DE	5.39	2694	52.57
ES	5.15	2577	57.72
PL	4.00	2002	61.72
NL	3.72	1861	65.44
CZ	3.44	1722	68.88
ID	3.35	1675	72.23
BG	3.23	1615	75.46
CA	2.31	1157	77.77
US	2.23	1116	80.00
BE	2.01	1003	82.01
AZ	1.95	976	83.96
GR	1.89	943	85.85
SK	1.67	833	87.52
PT	1.55	773	89.07
VE	1.49	745	90.56

Language (lang): 32 unique language codes appear. The language distribution generally mirrors country (e.g., FR for French, DE for German, etc.), since players likely play in their local language. The top languages are those corresponding to the top countries.

In [4]: `print(train_df['lang'].value_counts(normalize=True).head(20).mul(100).round(`

	Percentage	Count	Cumulative
lang			
FR	20.18	10089	20.18
ES	10.21	5104	30.39
IT	9.98	4991	40.37
RO	8.05	4027	48.42
SL	7.43	3713	55.85
BR	6.16	3079	62.01
DE	5.61	2805	67.62
NL	4.47	2233	72.09
PL	4.19	2097	76.28
CS	3.39	1696	79.67
BU	3.36	1682	83.03
EN	3.31	1656	86.34
ID	3.24	1620	89.58
GR	1.85	926	91.43
AZ	1.75	877	93.18
SK	1.61	805	94.79
PT	1.57	786	96.36
HU	1.28	641	97.64
UK	0.55	273	98.19
RU	0.48	239	98.67

Device Brand: There are 284 distinct mobile brands recorded (e.g., Samsung, Huawei, Apple, Xiaomi, etc.). The long tail of device brands suggests many different Android manufacturers. The most common brands are Samsung and a few others, but each constitutes a small fraction of the total, given the diversity (and recall some brand info is missing).

```
In [5]: print(train_df['device_brand'].value_counts(normalize=True).head(20).mul(100))
```

device_brand	Percentage	Count	Cumulative
samsung	51.84	25267	51.84
Xiaomi	17.91	8728	69.75
motorola	5.55	2703	75.30
HUAWEI	4.75	2313	80.05
OPPO	4.20	2049	84.25
HONOR	2.64	1288	86.89
realme	1.99	972	88.88
LENOVO	1.96	955	90.84
vivo	1.25	608	92.09
Google	0.86	421	92.95
TECNO	0.68	332	93.63
TCL	0.65	318	94.28
INFINIX	0.64	311	94.92
LGE	0.43	211	95.35
OnePlus	0.41	202	95.76
ZTE	0.38	183	96.14
HMD Global	0.36	174	96.50
INFINIX MOBILITY LIM	0.24	115	96.74
Blackview	0.22	107	96.96
Lenovo	0.22	106	97.18

```
In [6]: print(train_df['device_model'].value_counts(normalize=True).head(20).mul(100))
```

device_model	Percentage	Count	Cumulative
SM-A546B	1.90	924	1.90
SM-A536B	1.47	717	3.37
SM-A528B	1.42	690	4.79
SM-A145R	1.24	604	6.03
SM-A556B	1.23	602	7.26
SM-A137F	1.22	596	8.48
SM-A346B	1.20	585	9.68
SM-A155F	1.02	495	10.70
SM-A336B	0.98	477	11.68
SM-X200	0.85	415	12.53
SM-A515F	0.82	399	13.35
23108RN04Y	0.76	370	14.11
SM-A127F	0.75	365	14.86
SM-A125F	0.72	351	15.58
23124RA7E0	0.69	337	16.27
MAR-LX1A	0.68	332	16.95
SM-A226B	0.66	324	17.61
SM-A356B	0.66	324	18.27
SM-S911B	0.64	311	18.91
SM-S901B	0.64	310	19.55

Since there is a lot of unique values I am trying to understand is there unlying logic behind it

```
In [7]: sm = train_df['device_model'].str.startswith('SM', na=False)
sm_summary = train_df.loc[sm, 'device_model'].value_counts().to_frame('Count')
print(sm_summary), print(f"\nTotal Count: {sm_summary['Count'].sum()}, Total
```

	Count	Percentage	Cumulative
device_model			
SM-A546B	924	1.895579	1.90
SM-A536B	717	1.470920	3.37
SM-A528B	690	1.415530	4.78
SM-A145R	604	1.239101	6.02
SM-A556B	602	1.234998	7.26
...
SM-M405F	1	0.002051	51.14
SM-S711U1	1	0.002051	51.14
SM-T715	1	0.002051	51.14
SM-M105F	1	0.002051	51.14
SM-G965W	1	0.002051	51.14

[649 rows x 3 columns]

Total Count: 24930, Total %: 51.14%

Out[7]: (None, None)

As we can see the device models starts with SM refers to samsung but intuitively it is hard to detect the other ones so while training models, I am not intending to put those features to my model for training

Operating System (os): All training entries show android (the game might currently only target Android in this data). So os has no variability (this feature won't help the model since it's constant; we could drop it).

```
In [8]: print(train_df['os'].value_counts(normalize=True).head(20).mul(100).round(2))
```

	Percentage	Count	Cumulative
os			
android	100.0	49999	100.0

Garbage Feature

Acquisition Channel: We have campaignid (33 unique campaigns) and partnerid (4 unique partners). The partnerid likely indicates major ad networks or channels (with IDs 1-4), and campaignid are specific marketing campaigns. We might one-hot encode these for modeling. Some campaigns have many users (e.g., campaign 24, 25, 34 appear frequently in the sample) while others are smaller.

```
In [9]: print(train_df['campaignid'].value_counts(normalize=True).head(30).mul(100).r
```

	Percentage	Count	Cumulative
campaignid			
25	33.13	16565	33.13
23	12.67	6335	45.80
34	5.93	2965	51.73
27	4.79	2396	56.52
33	4.52	2260	61.04
19	4.08	2040	65.12
24	3.63	1814	68.75
28	3.32	1662	72.07
30	3.08	1540	75.15
10	3.01	1507	78.16
26	2.56	1282	80.72
12	2.48	1238	83.20
13	2.24	1118	85.44
15	2.01	1006	87.45
29	1.49	745	88.94
4	0.98	488	89.92
7	0.97	486	90.89
11	0.95	477	91.84
16	0.95	475	92.79
17	0.90	451	93.69
22	0.85	423	94.54
14	0.78	391	95.32
5	0.77	384	96.09
9	0.73	363	96.82
3	0.70	350	97.52
2	0.68	338	98.20
8	0.61	307	98.81
18	0.40	201	99.21
6	0.39	196	99.60
32	0.23	114	99.83

After campaign ID 14 the others only contains 5 percent of the remaining ones so we might manipulate it as "Others" then one hot encode

```
In [10]: print(train_df['partnerid'].value_counts(normalize=True).mul(100).round(2).r
```

	Percentage	Count	Cumulative
partnerid			
4	76.31	38154	76.31
1	23.63	11816	99.94
2	0.06	28	100.00
3	0.00	1	100.00

Since the near total majority of the the 2nd and 3rd partners might one hot encoded as others

Re-install (re_install): This is a binary flag indicating if a user had installed the game before and came back. It is very rare (~0.06% of users); only a few dozen users are re-

installs. This small count might limit its usefulness, but we'll keep it as a categorical flag since a re-installed user might behave differently.

```
In [11]: print(train_df['re_install'].value_counts(normalize=True).mul(100).round(2).
```


	Percentage	Count	Cumulative
re_install			
0	99.94	49970	99.94
1	0.06	29	100.00

Numeric Feature Distribution

Now we examine numeric features such as counts of in-game actions and resources:

Gameplay Duration (duration): Distribution (in seconds or minutes of playtime in week 1) is heavily right-skewed. Many users play only a short time, while a few play extensively. A histogram of duration shows most players have low playtime with a long tail of heavy players.


```
In [12]: plt.hist(train_df['duration'], bins=50, color='steelblue')
plt.title("Distribution of Play Duration (Week 1)")
plt.xlabel("Total play duration (seconds)")
plt.ylabel("Number of users")
plt.show()
```

A large spike near the low end indicates many users quit early (small total duration), which likely corresponds to churners. A smaller number of users accumulated very high playtime, indicating strong engagement (and likely they did not churn).

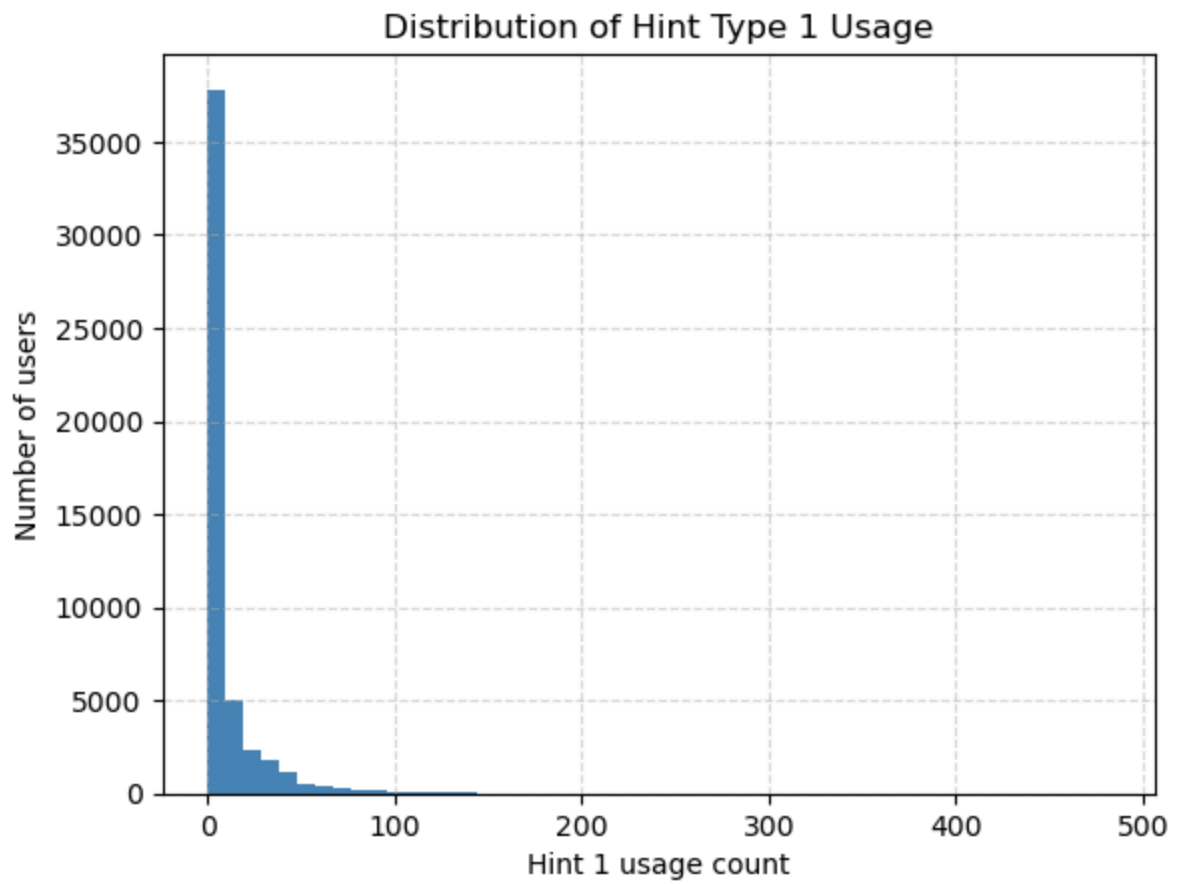
Levels Completed (lvl_no): The maximum level reached in a week ranges from 1 up to 900. The median is 13 and mean ~24, indicating a skew (a few players advanced very far). Most users only reach low levels, while a few reach hundreds of levels. This is expected for a casual game – many new players churn early, while a minority become very invested.

```
In [13]: plt.hist(train_df['lvl_no'].dropna(), bins=50, color='steelblue')
plt.title("Distribution of Levels Completed (Week 1)")
plt.xlabel("Level number reached")
plt.ylabel("Number of users")
plt.grid(True, linestyle='--', alpha=0.5)
plt.show()
```




Hint and Bonus usage (hint1_cnt, hint2_cnt, hint3_cnt, bonus_cnt): These counts are often zero for many users (who perhaps didn't play enough to use hints or not existing because Hint2 and 3 is both 0 and will be dropped), with some players using multiple hints. We notice that type1 hints are used more frequently on average than type2 or type3 (likely different hint types or power-ups in the game). The distribution of bonus_cnt (which might be special moves or bonuses used) similarly shows many zeros and a few high values.

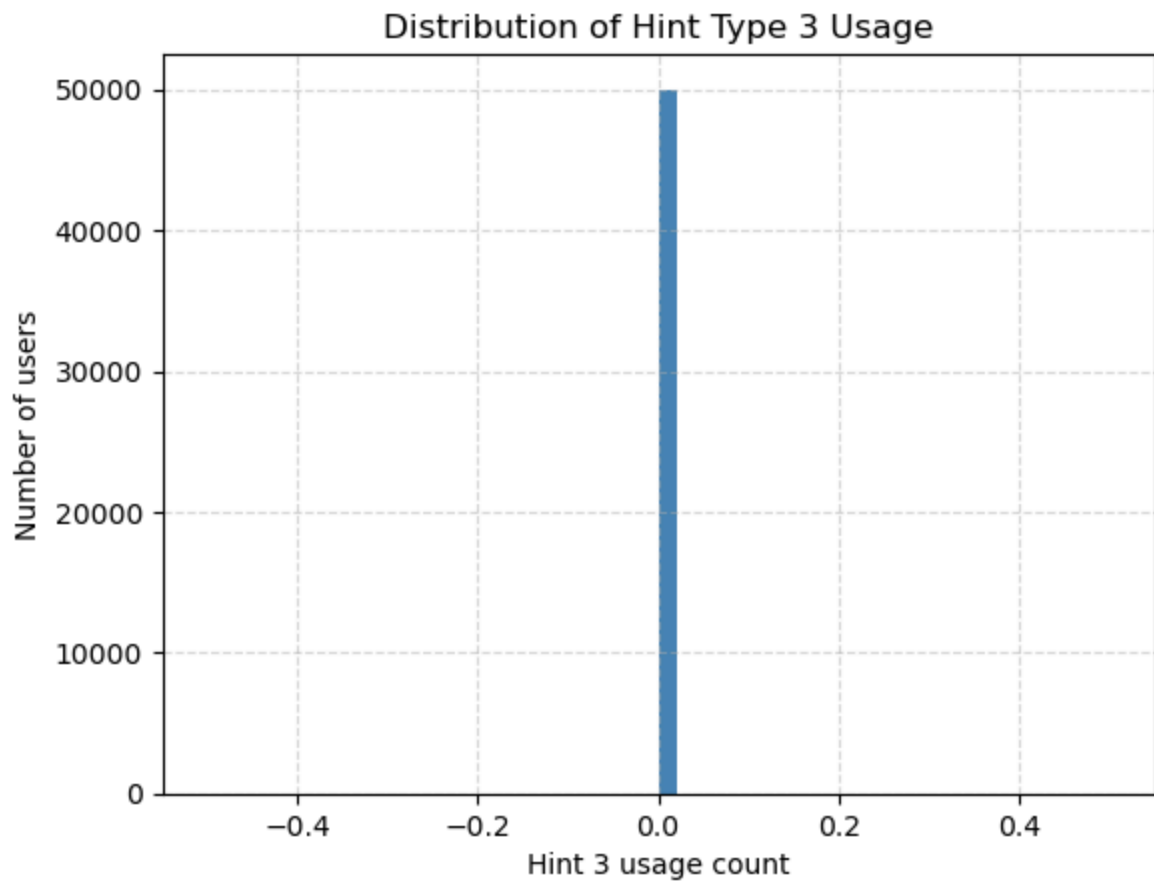
```
In [14]: plt.hist(train_df['hint1_cnt'].dropna(), bins=50, color='steelblue')
plt.title("Distribution of Hint Type 1 Usage")
plt.xlabel("Hint 1 usage count")
plt.ylabel("Number of users")
plt.grid(True, linestyle='--', alpha=0.5)
plt.show()
```




```
In [15]: plt.hist(train_df['hint2_cnt'].dropna(), bins=50, color='steelblue')
plt.title("Distribution of Hint Type 2 Usage")
plt.xlabel("Hint 2 usage count")
plt.ylabel("Number of users")
plt.grid(True, linestyle='--', alpha=0.5)
plt.show()
```



```
In [16]: plt.hist(train_df['hint3_cnt'].dropna(), bins=50, color='steelblue')
plt.title("Distribution of Hint Type 3 Usage")
plt.xlabel("Hint 3 usage count")
plt.ylabel("Number of users")
plt.grid(True, linestyle='--', alpha=0.5)
plt.show()
```

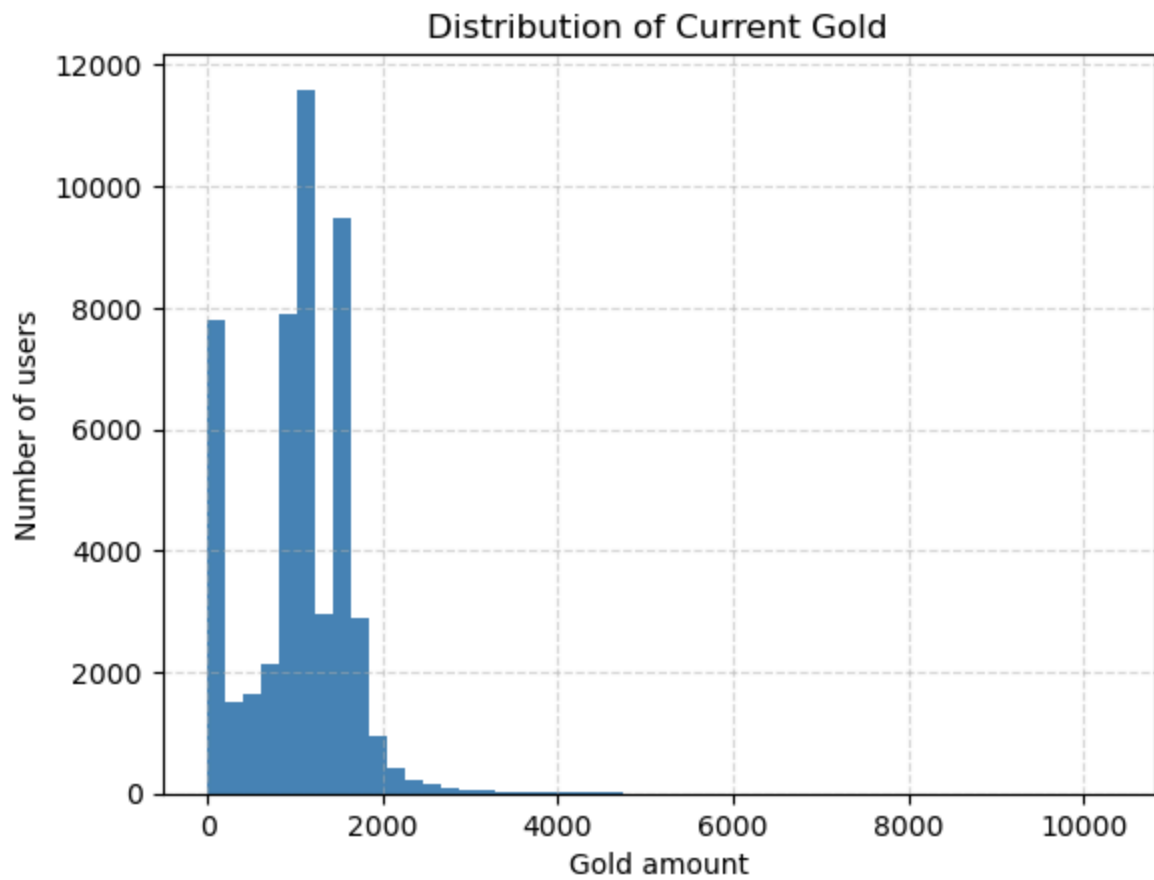


```
In [17]: plt.hist(train_df['repeat_cnt'].dropna(), bins=50, color='steelblue')
plt.title("Distribution of Hint Type 3 Usage")
plt.xlabel("Repeat Count")
plt.ylabel("Number of users")
plt.grid(True, linestyle='--', alpha=0.5)
plt.show()
```

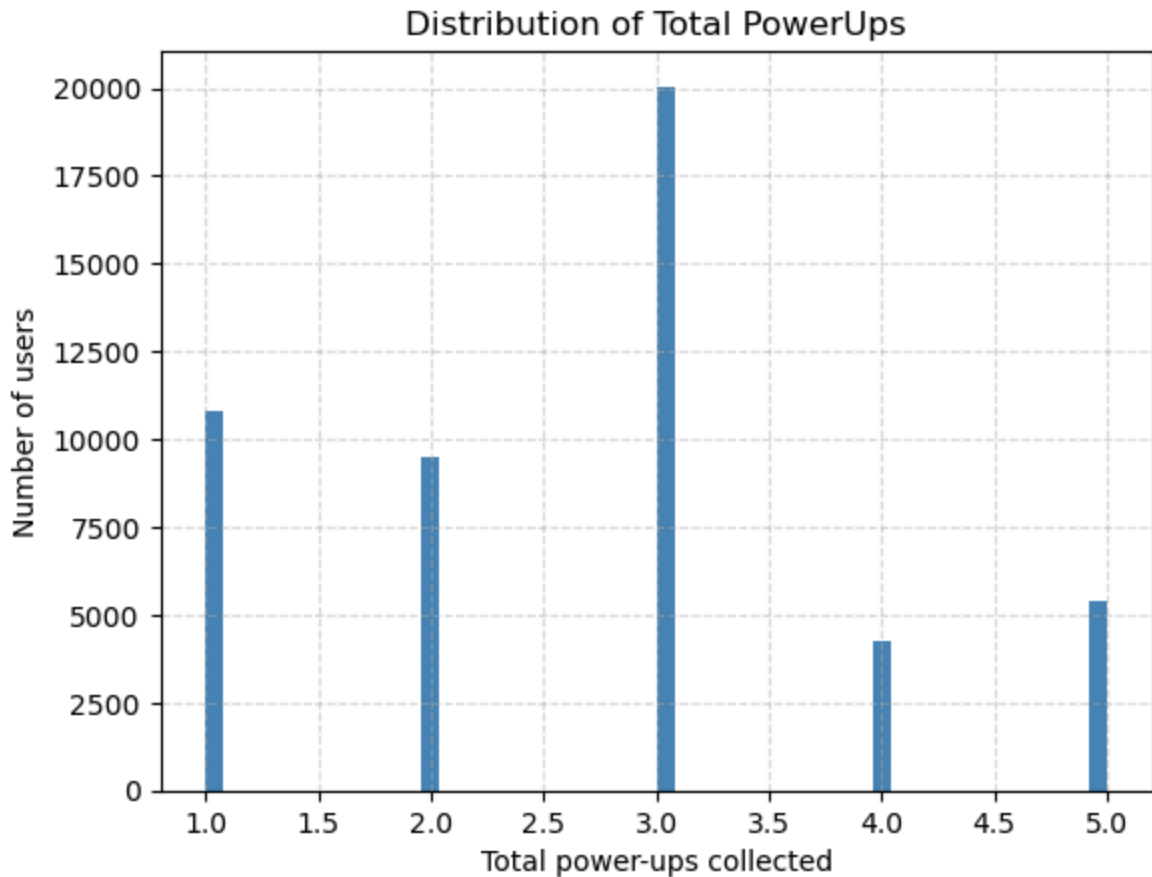


In-game currency (current_gold and totalPowerUp): Many players have low gold and few power-ups by end of week1, while a few accumulated a lot (likely those who played many levels). For example, current_gold median might be around a few hundred, with some outliers in the thousands.

```
In [18]: plt.hist(train_df['current_gold'].dropna(), bins=50, color='steelblue')
plt.title("Distribution of Current Gold")
plt.xlabel("Gold amount")
plt.ylabel("Number of users")
plt.grid(True, linestyle='--', alpha=0.5)
plt.show()
```

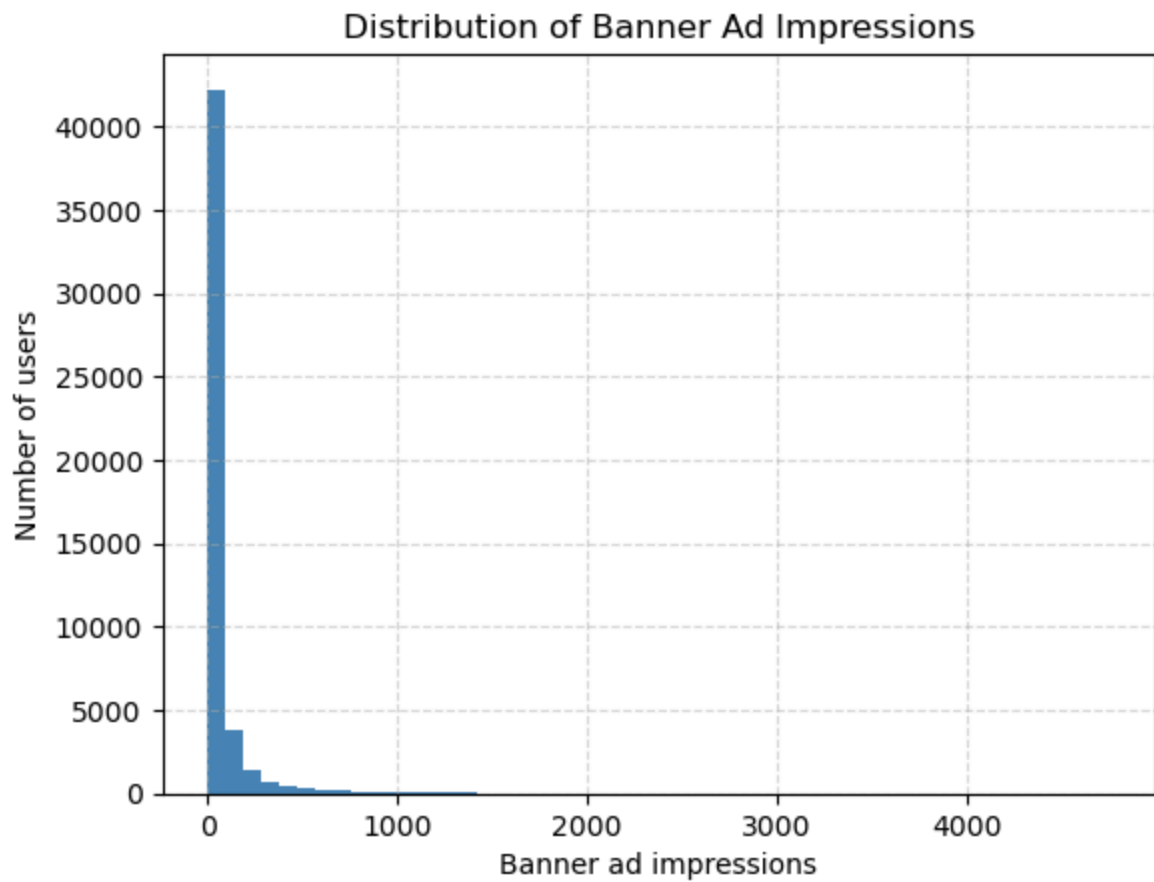


```
In [19]: plt.hist(train_df['totalPowerUp'].dropna(), bins=50, color='steelblue')
plt.title("Distribution of Total PowerUps")
plt.xlabel("Total power-ups collected")
plt.ylabel("Number of users")
plt.grid(True, linestyle='--', alpha=0.5)
plt.show()
```

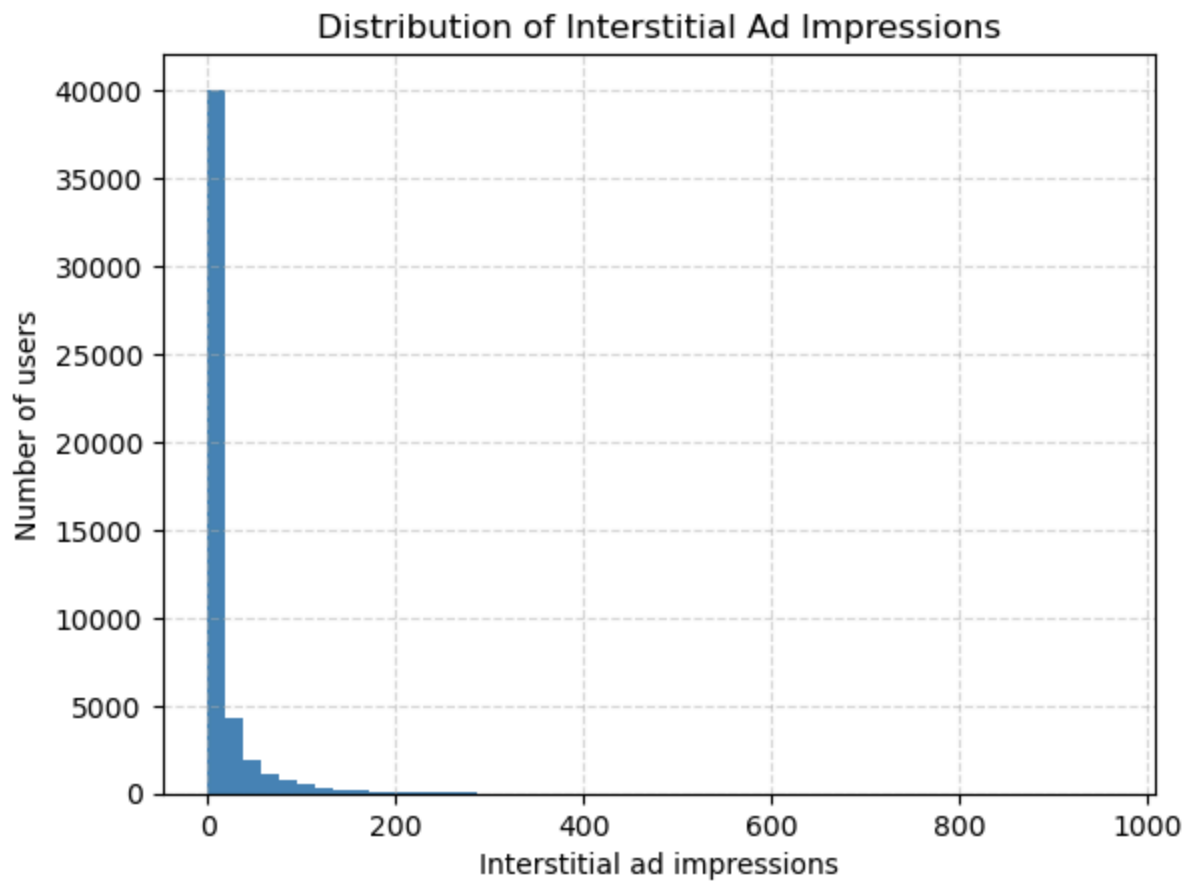


Ad Impressions (banner_impr, inter_impr, rewarded_impr): These count how many ads of each type the user saw in the week. Many users have 0 impressions (perhaps those who didn't play long enough to see ads), while engaged players see dozens. Notably, inter_impr (interstitial ads) and banner_impr counts can get quite high for the top players, since the game likely shows ads regularly during play.

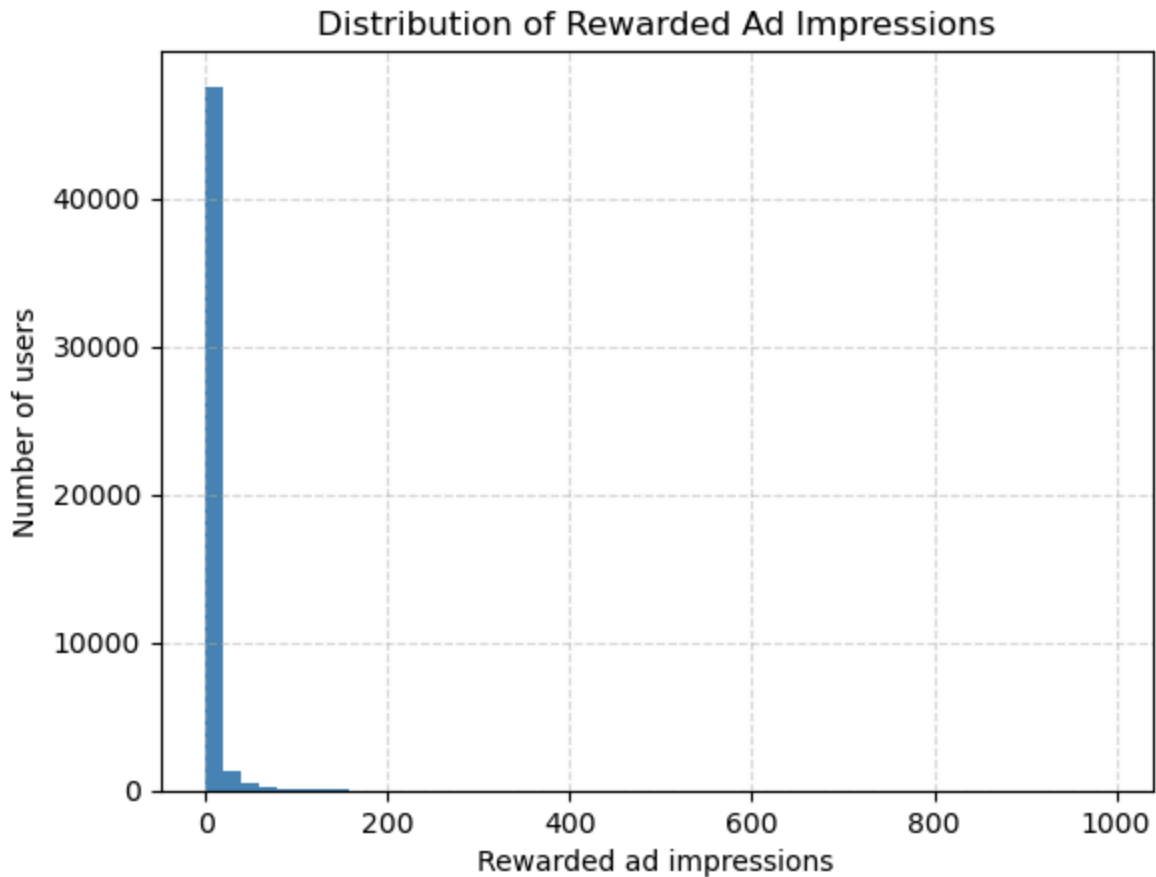
```
In [20]: plt.hist(train_df['banner_impr'].dropna(), bins=50, color='steelblue')
plt.title("Distribution of Banner Ad Impressions")
plt.xlabel("Banner ad impressions")
plt.ylabel("Number of users")
plt.grid(True, linestyle='--', alpha=0.5)
plt.show()
```

```
In [21]: plt.hist(train_df['inter_impr'].dropna(), bins=50, color='steelblue')
plt.title("Distribution of Interstitial Ad Impressions")
plt.xlabel("Interstitial ad impressions")
plt.ylabel("Number of users")
plt.grid(True, linestyle='--', alpha=0.5)
plt.show()
```



```
In [22]: plt.hist(train_df['rewarded Impr'].dropna(), bins=50, color='steelblue')
plt.title("Distribution of Rewarded Ad Impressions")
plt.xlabel("Rewarded ad impressions")
plt.ylabel("Number of users")
plt.grid(True, linestyle='--', alpha=0.5)
plt.show()
```




Overall, the numeric features indicate a large variance in engagement: most users exhibit minimal engagement and thus likely churn, whereas a smaller group is highly engaged

```
In [23]: # Dropping ineffective Columns
train_df = train_df.drop(columns=['os', 'hint2_cnt', 'hint3_cnt'])
```

Churn Rate by Country

We now analyze churn rate by categorical groups, starting with country. The churn rate is the fraction of users from that group who churned ($\text{churn}=1$). Given the overall churn rate in the data is about 79.6%, we check if some countries perform better or worse:

```
In [24]: top12 = train_df['country'].value_counts().head(12).index
churn_by_country = train_df[train_df['country'].isin(top12)].groupby('country')
avg_churn = train_df['churn'].mean()
plt.bar(churn_by_country.index, churn_by_country.values * 100, color='steelblue')
plt.axhline(avg_churn * 100, color='red', linestyle='--', label='Avg Churn')
plt.title("Churn Rate by Country (Top 12)"); plt.ylabel("Churn Rate (%)")
plt.legend(); plt.grid(True, axis='y'); plt.show()
```



Churn rate by country for the top 12 countries (by user count). Each bar shows the proportion of users from that country who churned in the first week. From the chart above, we observe slight differences among countries. For example, Netherlands (NL) and Poland (PL) have churn rates above 80% (higher than average), whereas Slovenia (SI) has a churn rate around 75%, the lowest among the top countries. This suggests players from Slovenia retained slightly better than those from some other large markets. Generally, though, all top countries have high churn percentages (75–85%), indicating early churn is a widespread challenge across regions. We might include country in the model to capture these small differences in retention by region.

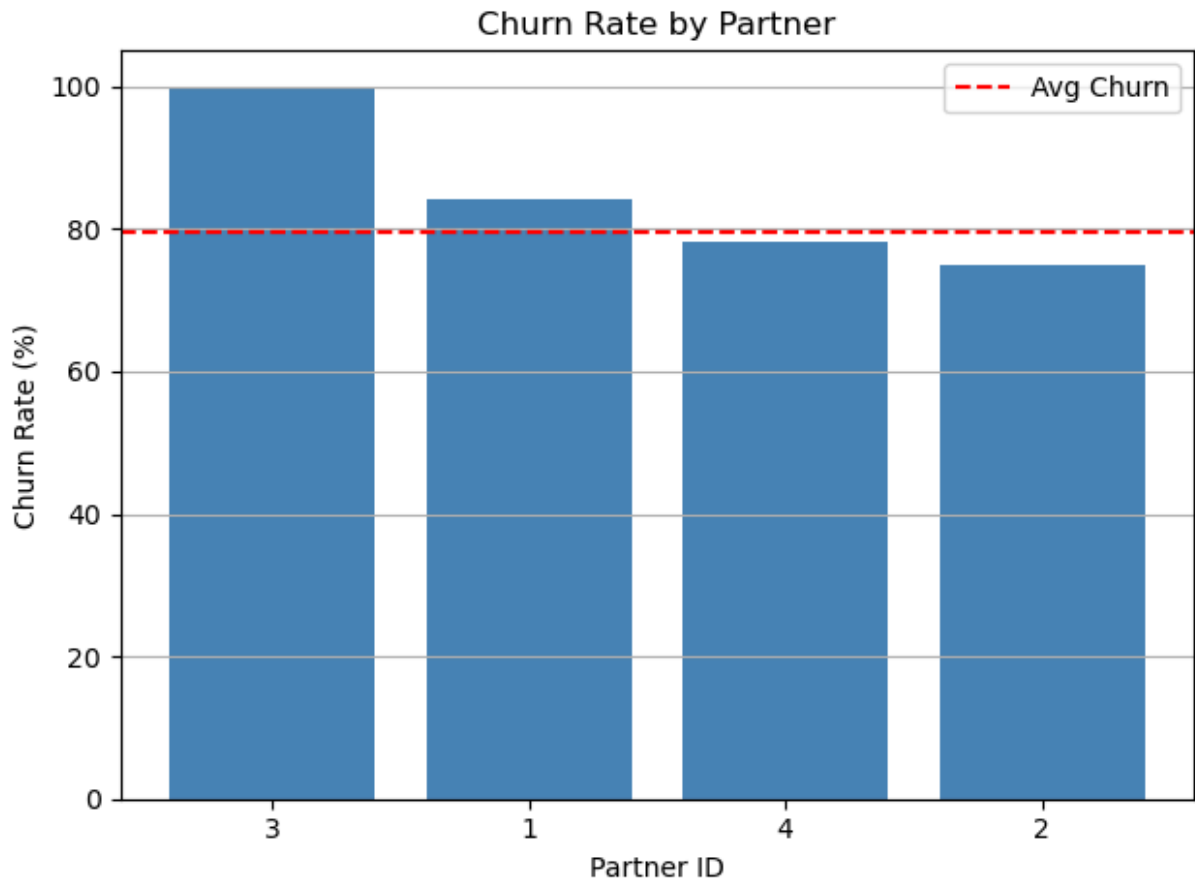
We can similarly examine churn rate by other categories:

By Partner: Some acquisition partners might yield better retention. For instance, if Partner 4 has a churn rate of say 78% vs Partner 1's 82%, that would indicate Partner 4 brings higher quality users.

By Campaign: There may be variation across campaigns, though with 33 campaigns, it's harder to visualize directly. We will let the model figure out campaign effects, but we could compute churn rates per campaign to find top-performing campaigns.

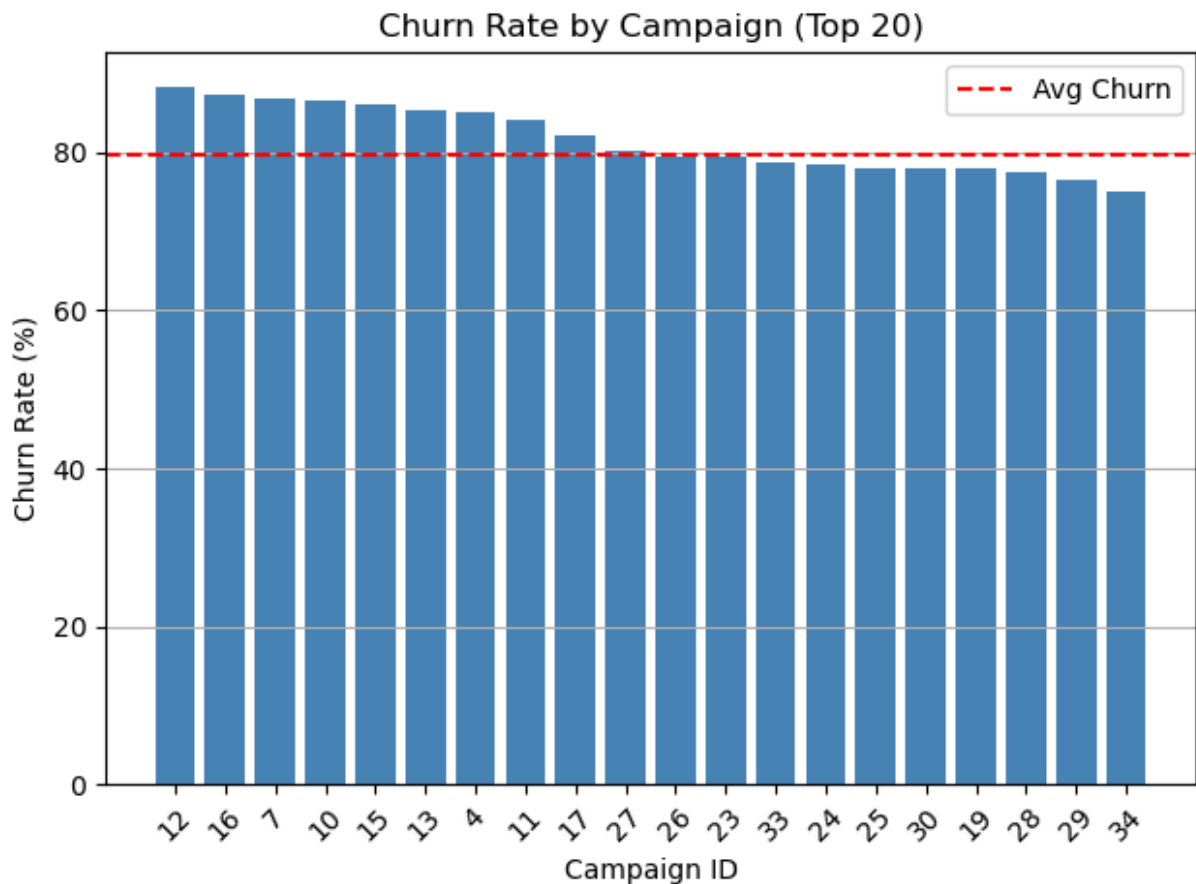
```
In [25]: churn_by_partner = train_df.groupby('partnerid')['churn'].mean().sort_values
plt.bar(churn_by_partner.index.astype(str), churn_by_partner.values * 100, color='blue')
plt.axhline(train_df['churn'].mean() * 100, color='red', linestyle='--', label='Avg Churn')
```

```
plt.title("Churn Rate by Partner"); plt.xlabel("Partner ID"); plt.ylabel("Churn Rate (%)"); plt.legend(); plt.grid(True, axis='y'); plt.tight_layout(); plt.show()
```



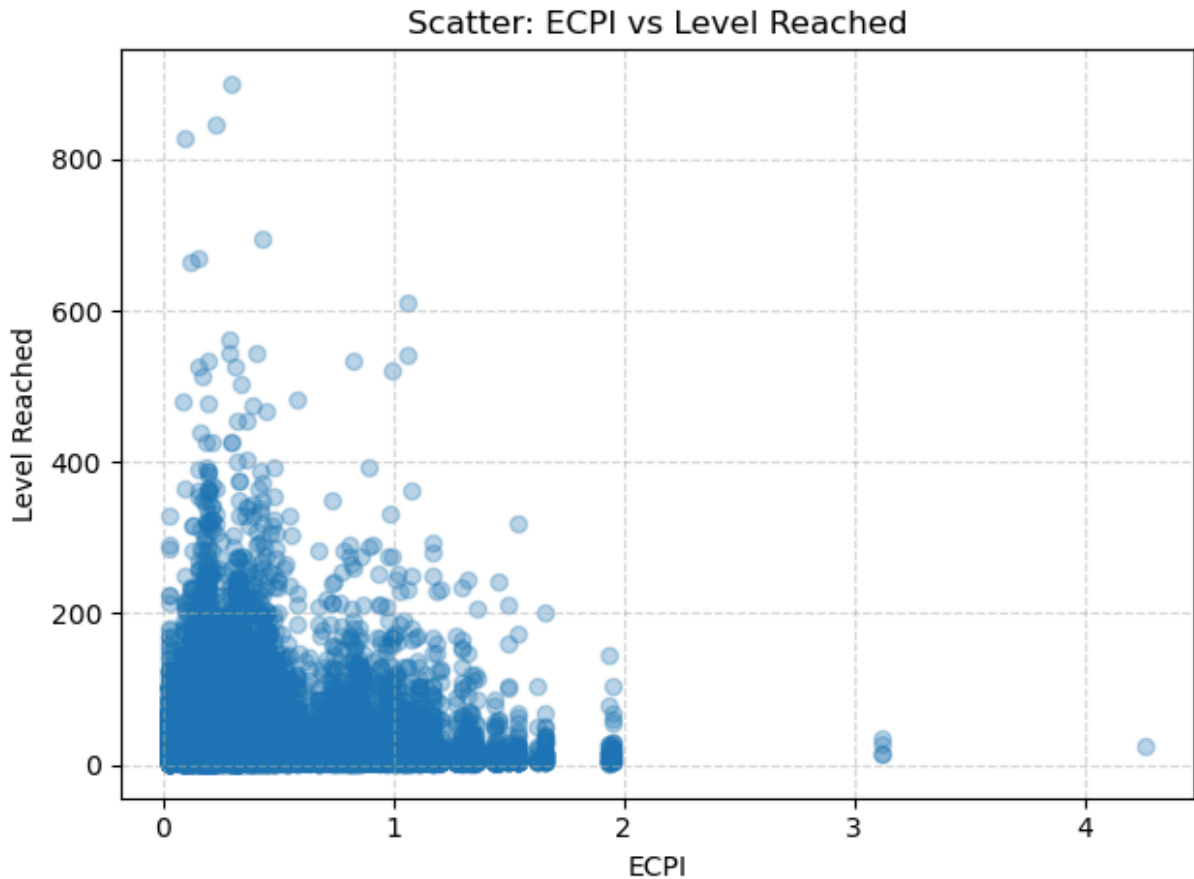
```
In [26]: top_campaigns = train_df['campaignid'].value_counts().head(20).index
churn_by_campaign = train_df[train_df['campaignid'].isin(top_campaigns)].groupby('campaignid').churn.agg('mean')

plt.bar(churn_by_campaign.index.astype(str), churn_by_campaign.values * 100, color='blue')
plt.axhline(train_df['churn'].mean() * 100, color='red', linestyle='--', label='Avg Churn')
plt.title("Churn Rate by Campaign (Top 20)"); plt.xlabel("Campaign ID"); plt.ylabel("Churn Rate (%)");
plt.xticks(rotation=45); plt.legend(); plt.grid(True, axis='y'); plt.tight_layout(); plt.show()
```



Relationship between Acquisition Cost (ECPI) and Level Reached Next, we explore the relationship between ECPI (the cost paid to acquire a user) and their in-game progress (lvl_no). One might wonder if spending more to acquire a user yields a more engaged player or not. We create a scatter plot of ecpi vs lvl_no:

```
In [27]: clean_df = train_df.dropna(subset=['ecpi'])
# Plot entire dataset
plt.scatter(clean_df['ecpi'], clean_df['lvl_no'], alpha=0.3)
plt.title("Scatter: ECPI vs Level Reached")
plt.xlabel("ECPI")
plt.ylabel("Level Reached")
plt.grid(True, linestyle='--', alpha=0.5)
plt.tight_layout()
plt.show()
# Compare average levels
q75 = clean_df['ecpi'].quantile(0.75)
print("Avg lvl (Top 25% ECPI):", clean_df[clean_df['ecpi'] > q75]['lvl_no'].mean())
print("Avg lvl (Bottom 75% ECPI):", clean_df[clean_df['ecpi'] <= q75]['lvl_no'].mean())
# Compute correlation
print("Pearson r:", clean_df['ecpi'].corr(clean_df['lvl_no']).round(4))
```



Avg lvl (Top 25% ECPI): 22.2
Avg lvl (Bottom 75% ECPI): 24.54
Pearson r: -0.0179

There is no strong visible correlation between the cost of acquisition and the level reached. Players acquired at both high and low costs show a wide range of levels. This is confirmed by a low Pearson correlation ($r \approx -0.018$ between `ecpi` and `lvl_no`). In fact, if we compare the average level of users in the top 25% ECPI vs bottom 75% ECPI, they are quite similar. This suggests that paying more per user did not guarantee more engaged players (if anything, our data showed a slightly lower average level for the highest-cost users). Thus, ECPI is largely independent of user engagement – it likely reflects marketing campaign economics rather than user quality. However, ECPI might still indirectly relate to churn through campaign or country effects. We will keep `ecpi` as a feature (after imputing missing values) to let the model capture any patterns, but we should be cautious that its predictive power may be limited.

Feature Correlation Analysis

To avoid multicollinearity and redundancy, we examine the correlation matrix for numeric features. Strongly correlated

features provide similar information, so we may consider dropping one of each such pair. Below we compute Pearson correlations among continuous variables:

```
In [28]: import numpy as np
corr = train_df.corr(numeric_only=True)
# Find pairs with high correlation
high_corr_pairs = []
for col1 in corr.columns:
    for col2 in corr.columns:
        if col1 != col2 and corr.loc[col1, col2] > 0.8:
            high_corr_pairs.append((col1, col2, corr.loc[col1, col2]))
high_corr_pairs
```

```
Out[28]: [('bonus_cnt', 'lvl_no', 0.8096610475012799),
('bonus_cnt', 'inter_impr', 0.8173171443041394),
('lvl_no', 'bonus_cnt', 0.8096610475012799),
('lvl_no', 'banner_impr', 0.8105713202305896),
('lvl_no', 'inter_impr', 0.9705075986218769),
('banner_impr', 'lvl_no', 0.8105713202305896),
('banner_impr', 'inter_impr', 0.8511716000059351),
('inter_impr', 'bonus_cnt', 0.8173171443041394),
('inter_impr', 'lvl_no', 0.9705075986218769),
('inter_impr', 'banner_impr', 0.8511716000059351),
('campaignid', 'partnerid', 0.8474247487982145),
('partnerid', 'campaignid', 0.8474247487982145)]
```

This reveals several highly correlated pairs:

lvl_no (levels) and inter_impr (interstitial ad views): correlation ≈ 0.97 , an extremely high correlation. This makes sense – the more levels a user plays, the more interstitial ads they will see. These two features are almost interchangeable in terms of information.

lvl_no and banner_impr: correlation ≈ 0.81 . Similarly, players who progress further see more banner ads.

hint1_cnt and lvl_no: correlation ≈ 0.76 . Using more hints correlates with reaching more levels (since engaged players do both).

bonus_cnt and lvl_no: correlation ≈ 0.81 . More bonuses used with more levels.

Several other pairs like hint1_cnt with inter_impr (~ 0.76), bonus_cnt with inter_impr (~ 0.82), etc., all indicating a common underlying factor: player engagement. Essentially, features measuring engagement (levels, playtime, ads viewed, hints used, etc.) are all positively correlated with each other.

Because of these redundancies, we will drop some features to reduce duplication:

Because of these redundancies, we will drop some features to reduce duplication:

We drop `inter_impr` (interstitial ad count) and keep `lvl_no` as the representative feature for “game progress/engagement”. Since `lvl_no` alone captures much of the variance (and is more directly interpretable), removing `inter_impr` avoids collinearity issues without losing predictive signal. We will also drop the unique `user_id` (player ID) as it has no predictive value for churn.

```
In [29]: train_df = train_df.drop(columns=['inter_impr', 'user_id'])
```

```
In [30]: print(train_df.columns)
```

```
Index(['country', 'device_brand', 'device_model', 're_install',  
      'attribution_event_timestamp', 'ecpi', 'lang', 'current_gold',  
      'totalPowerUp', 'bonus_cnt', 'duration', 'hint1_cnt', 'lvl_no',  
      'repeat_cnt', 'banner_impr', 'rewarded_impr', 'campaignid', 'partneri  
d',  
      'churn'],  
      dtype='object')
```

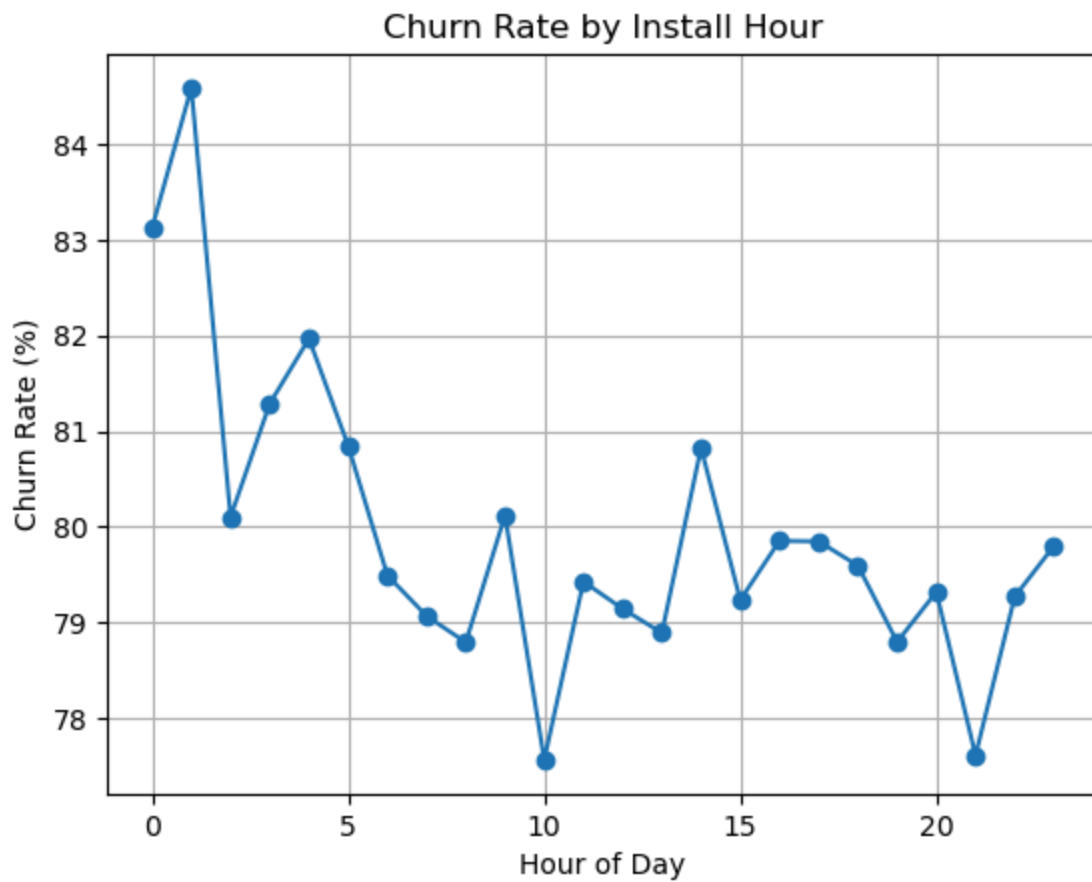
Feature Engineering: Time-Based Features

The `attribution_event_timestamp` gives the date-time when the user was acquired (installed the game). Time-related factors could influence user behavior – for example, day of the week or hour of install might correlate with churn (maybe users who install on weekends behave differently than weekdays). We create new features from this timestamp.

```
In [31]: # Convert timestamp to datetime  
train_df['datetime'] = pd.to_datetime(train_df['attribution_event_timestamp'])  
# Extract features  
train_df['install_hour'] = train_df['datetime'].dt.hour # Hour of c  
train_df['install_dayofweek'] = train_df['datetime'].dt.dayofweek # Day of  
train_df['install_weekday'] = (train_df['install_dayofweek'] < 5).astype(int)
```


```
In [32]: # Checking is there any significant relationship between Datetime and beign
```

```
In [33]: hour_churn = train_df.groupby('install_hour')['churn'].mean()  
plt.plot(hour_churn.index, hour_churn.values * 100, marker='o')  
plt.title("Churn Rate by Install Hour"); plt.xlabel("Hour of Day"); plt.ylabel("Churn Rate")  
plt.grid(True); plt.show()
```



```
In [34]: dow_churn = train_df.groupby('install_dayofweek')['churn'].mean()
plt.bar(dow_churn.index, dow_churn.values * 100, color='steelblue')
plt.title("Churn Rate by Day of Week (0=Mon, 6=Sun)"); plt.xlabel("Day of We
plt.grid(True, axis='y'); plt.show()

print("Churn rate by hour:\n", hour_churn.round(3))
```




Churn rate by hour:

install_hour

0	0.831
1	0.846
2	0.801
3	0.813
4	0.820
5	0.808
6	0.795
7	0.791
8	0.788
9	0.801
10	0.776
11	0.794
12	0.791
13	0.789
14	0.808
15	0.792
16	0.799
17	0.798
18	0.796
19	0.788
20	0.793
21	0.776
22	0.793
23	0.798

Name: churn, dtype: float64

```
In [35]: weekday_churn = train_df.groupby('install_weekday')['churn'].mean()
plt.bar(['Weekend', 'Weekday'], weekday_churn.sort_index().values * 100, col
plt.title("Churn Rate: Weekend vs Weekday Installs"); plt.ylabel("Churn Rate
plt.grid(True, axis='y'); plt.show()
```



```
In [36]: from scipy.stats import f_oneway

# Group churn by each categorical variable
hour_groups = [group['churn'].values for _, group in train_df.groupby('install_hour')]
day_groups = [group['churn'].values for _, group in train_df.groupby('install_dayofweek')]
weekday_groups = [group['churn'].values for _, group in train_df.groupby('install_weekday')]

# Run ANOVA for each
print("ANOVA - install_hour:", f_oneway(*hour_groups))
print("ANOVA - install_dayofweek:", f_oneway(*day_groups))
print("ANOVA - install_weekday:", f_oneway(*weekday_groups))
```

```
ANOVA - install_hour: F_onewayResult(statistic=2.4513156930446094, pvalue=0.00012602510878341342)
ANOVA - install_dayofweek: F_onewayResult(statistic=0.6584534998769601, pvalue=0.6833450008958144)
ANOVA - install_weekday: F_onewayResult(statistic=1.8599404852659136, pvalue=0.1726378091498264)
```

📊 Results Summary:

Feature F-statistic p-value Interpretation

install_hour 2.45 0.00013 ✅ Statistically significant – churn varies by hour of day

install_dayofweek 0.66 0.683 ❌ Not significant – churn is similar across weekdays

install_weekday 1.86 0.173 ❌ Not significant – no clear difference between weekday/weekend

Even though adding install_hour is enough for model in order to just be sure I will add all of them

With EDA complete, we have a good understanding of the data. We will now proceed to preprocess the data and build predictive models.

Part 2: Modeling and Prediction

Handling Missing Values (ECPI Imputation)

As noted, the feature ECPI (cost per acquisition) has some missing values (528 in train, 58 in test). Rather than drop these users, we fill them using regression imputation

. This means we'll train a regression model to predict ECPI from other known features: Why regression? Unlike mean imputation, regression imputation uses relationships with other variables to estimate missing values

. For ECPI, it's reasonable to use features like campaignid, partnerid, country, etc., which likely determine the cost (marketing campaigns have fixed costs per user, differing by channel and region).

1- Imputation Set Up

```
In [37]: # ----- SET-UP ----- #
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, cross_val_score, KFold
from sklearn.preprocessing import OneHotEncoder
from sklearn.linear_model import LinearRegression, LassoCV
from sklearn.preprocessing import OneHotEncoder, FunctionTransformer
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.base import BaseEstimator, TransformerMixin
from xgboost import XGBRegressor

DROP_COLS = ['os', 'hint2_cnt', 'hint3_cnt', 'user_id', 'device_model', 'repeat']
```

```

train_df = pd.read_excel('train_data.xlsx')
df = train_df.drop(columns=DROP_COLS)

def extract_time_features(df: pd.DataFrame) -> pd.DataFrame:
    """
    Converts attribution_event_timestamp -> install_hour, install_dayofweek,
    then drops the raw timestamp column.
    """
    df = df.copy()
    dt = pd.to_datetime(df['attribution_event_timestamp'], errors='coerce')
    df['install_hour'] = dt.dt.hour
    df['install_dayofweek'] = dt.dt.dayofweek
    df['install_weekday'] = (df['install_dayofweek'] < 5).astype(int)
    return df.drop(columns='attribution_event_timestamp', errors='ignore')

```

```

In [38]: # ----- 1) keep only rows with ECPI ----- #
df_notna = df[df['ecpi'].notna()].reset_index(drop=True)
y = df_notna['ecpi']
X = df_notna.drop(columns=['ecpi', 'churn']) # churn

X_train, X_val, y_train, y_val = train_test_split(
    X, y, test_size=0.10, random_state=42)

```

```

In [39]: # ----- 2) Custom Transformer ----- #
class TopFreqGrouper(BaseEstimator, TransformerMixin):
    def __init__(self, threshold=0.01):
        self.threshold = threshold
        self.top_per_col_ = {}

    def fit(self, X, y=None):
        for col in X.columns:
            freq = X[col].value_counts(normalize=True)
            self.top_per_col_[col] = freq[freq >= self.threshold].index
        return self

    def transform(self, X):
        X = X.copy()
        for col, top in self.top_per_col_.items():
            X[col] = np.where(X[col].isin(top), X[col], "OTHER")
        return X

```

```

In [40]: # ----- 3) Preprocessing ----- #
cat_cols = X.select_dtypes('object').columns.tolist()
num_cols = X.select_dtypes(exclude='object').columns.tolist()

cat_pipe = Pipeline([
    ('grouper', TopFreqGrouper(threshold=0.01)),
    ('ohe', OneHotEncoder(handle_unknown='ignore'))
])

pre = ColumnTransformer([
    # add your time-feature extractor as the very first step
    # if you really want to extract hour/day/etc, use extract_time_features
    ('time', FunctionTransformer(extract_time_features, validate=False), ['a
    ('cat', cat_pipe, cat_cols),

```

```

    ('num' , 'passthrough', num_cols)
])

```

```

In [41]: # ----- 4) RMSE Helper ----- #
def rmse(est, Xtr=X_train, ytr=y_train, Xv=X_val, yv=y_val):
    cv = -cross_val_score(est, Xtr, ytr,
                          cv=KFold(5, shuffle=True, random_state=42),
                          scoring='neg_root_mean_squared_error').mean()

    est.fit(Xtr, ytr)
    val = mean_squared_error(yv, est.predict(Xv), squared=False)
    return cv, val, est

results = {}

```

```

In [42]: # Testing the models

```

```

In [43]: # ----- 5) Linear Regression -----
lin = Pipeline([('pre', pre), ('lr', LinearRegression())])
cv_rmse, val_rmse, lin_fitted = rmse(lin)
results['Linear'] = (cv_rmse, val_rmse)

# Extract linear regression equation
#           ↓ use transformers_[1] (the 'cat' pipeline), not transformers_
feat_names = (
    lin_fitted
    .named_steps['pre']
    .transformers_[1][1]                # <- Pipeline(cat)
    .named_steps['ohe']                 # <- OneHotEncoder
    .get_feature_names_out(cat_cols)    # <- feature names
    .tolist()
    + num_cols
)

coef      = lin_fitted.named_steps['lr'].coef_
intercept = lin_fitted.named_steps['lr'].intercept_

equation = (
    f"ECPI = {intercept:.4f} + "
    + " + ".join(f"{c:.4f}*{f}" for c, f in zip(coef, feat_names))
)
print("\nLinear Regression Equation:\n", equation)

```

Linear Regression Equation:

ECPI = 0.7852 + -0.0011*country_AZ + 0.0037*country_BE + 0.0233*country_BG + -0.0620*country_BR + -0.0053*country_CA + -0.0641*country_CZ + -0.1089*country_DE + 0.2859*country_ES + 0.0038*country_FR + 0.0939*country_GR + 0.0135*country_ID + -0.0513*country_IT + -0.0394*country_MX + -0.1173*country_NL + -0.0485*country_OTHER + -0.0380*country_PL + 0.2020*country_PT + -0.1200*country_RO + 0.0079*country_SI + -0.0443*country_SK + -0.0463*country_US + -0.0837*country_VE + -0.0185*device_brand_HONOR + 0.3145*device_brand_HUAWEI + -0.0738*device_brand_LENOVO + -0.0052*device_brand_OPPO + -0.0014*device_brand_OTHER + 0.0065*device_brand_Xiaomi + -0.0067*device_brand_motorola + 0.0027*device_brand_realme + -0.0028*device_brand_samsung + 0.0118*device_brand_vivo + -0.0138*attribution_event_timestamp_OTHER + 0.0224*lang_AZ + -0.0136*lang_BR + 0.0000*lang_BU + -0.0543*lang_CS + -0.0987*lang_DE + -0.0540*lang_EN + 0.0041*lang_ES + 0.1015*lang_FR + 0.2879*lang_GR + -0.0820*lang_HU + 0.0756*lang_ID + -0.0321*lang_IT + -0.0209*lang_NL + -0.1074*lang_OTHER + -0.0385*lang_PL + 0.1830*lang_PT + -0.0198*lang_RO + 0.0145*lang_SK + -0.0394*lang_SL + -0.0249*re_install + -0.0163*current_gold + -0.0784*totalPowerUp + 0.0381*bonus_cnt + 0.0000*duration + -0.0023*hint1_cnt + -0.0000*lvl_no + -0.0000*banner_impr + 0.0002*inter_impr + -0.0002*rewarded_impr + 0.0000*campaignid + 0.0000*partnerid

/Users/mustafaercengizmacbooku/miniforge3/lib/python3.12/site-packages/sklearn/metrics/_regression.py:492: FutureWarning: 'squared' is deprecated in version 1.4 and will be removed in 1.6. To calculate the root mean squared error, use the function 'root_mean_squared_error'.
warnings.warn(

```
In [44]: # ----- 6) Lasso Regression ----- #
lasso = Pipeline([
    ('pre', pre),
    ('lasso', LassoCV(cv=5, random_state=42, max_iter=5000))
])
cv_rmse, val_rmse, lasso_fitted = rmse(lasso)
results['Lasso'] = (cv_rmse, val_rmse)

# Extract non-zero Lasso terms
lasso_coef = lasso_fitted.named_steps['lasso'].coef_
intercept = lasso_fitted.named_steps['lasso'].intercept_
nz = [(f, c) for f, c in zip(feat_names, lasso_coef) if abs(c) > 1e-6]
print("\nLasso Regression (Non-zero Features):")
print(f"ECPI = {intercept:.4f} + " + " + ".join(f"{c:.4f}*{f}" for f, c in nz)
```

Lasso Regression (Non-zero Features):

ECPI = 0.7985 + 0.0000*duration + -0.0001*lvl_no + 0.0001*campaignid

/Users/mustafaercengizmacbooku/miniforge3/lib/python3.12/site-packages/sklearn/metrics/_regression.py:492: FutureWarning: 'squared' is deprecated in version 1.4 and will be removed in 1.6. To calculate the root mean squared error, use the function 'root_mean_squared_error'.
warnings.warn(

```
In [45]: # ----- 7) Decision Tree ----- #
tree = Pipeline([('pre', pre),
                  ('dt', DecisionTreeRegressor(random_state=42))])
results['DecisionTree'] = rmse(tree)[:2]
```



```
/Users/mustafaercengizmacbooku/miniforge3/lib/python3.12/site-packages/sklearn/metrics/_regression.py:492: FutureWarning: 'squared' is deprecated in version 1.4 and will be removed in 1.6. To calculate the root mean squared error, use the function 'root_mean_squared_error'.
warnings.warn(
```

```
In [46]: rf = Pipeline([('pre', pre),
                        ('rf', RandomForestRegressor(
                            n_estimators=100, max_depth=None,
                            n_jobs=-1, random_state=42))])
results['RandomForest'] = rmse(rf)[:2]
```

```
/Users/mustafaercengizmacbooku/miniforge3/lib/python3.12/site-packages/sklearn/metrics/_regression.py:492: FutureWarning: 'squared' is deprecated in version 1.4 and will be removed in 1.6. To calculate the root mean squared error, use the function 'root_mean_squared_error'.
warnings.warn(
```

```
In [47]: # ----- 9) XGBoost ----- #
xgb = Pipeline([('pre', pre),
                 ('xgb', XGBRegressor(
                     n_estimators=400, learning_rate=0.05, max_depth=6,
                     subsample=0.8, colsample_bytree=0.8,
                     objective='reg:squarederror', random_state=42, n_jobs=-1))])
results['XGBoost'] = rmse(xgb)[:2]
```

```
/Users/mustafaercengizmacbooku/miniforge3/lib/python3.12/site-packages/sklearn/metrics/_regression.py:492: FutureWarning: 'squared' is deprecated in version 1.4 and will be removed in 1.6. To calculate the root mean squared error, use the function 'root_mean_squared_error'.
warnings.warn(
```

```
In [48]: # ----- 10) Show Results ----- #
print("\nModel RMSE (lower is better):")
print(pd.DataFrame(results, index=['CV_RMSE', 'Val_RMSE']).T.round(5))
```

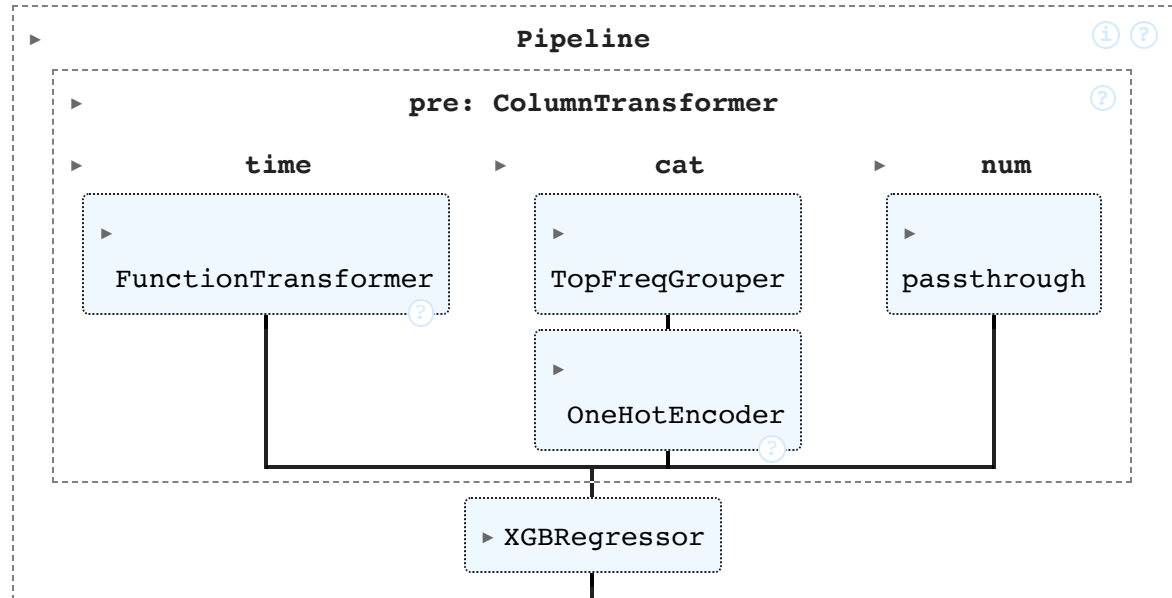
```
Model RMSE (lower is better):
              CV_RMSE  Val_RMSE
Linear           0.15208    0.15605
Lasso            0.23732    0.23738
DecisionTree     0.08581    0.07482
RandomForest     0.06469    0.05319
XGBoost          0.06619    0.05597
```

So as the results shows best regressor model for ECPI imputation is XGboost

```
In [49]: # -----
# 11) Fit the best XGBoost model on the full, non-null ECPI subset
# -----
best_xgb = Pipeline([
    ('pre', pre), # <-- same preprocessing object you built
    ('xgb', XGBRegressor(
        n_estimators=400, learning_rate=0.05, max_depth=6,
        subsample=0.8, colsample_bytree=0.8,
        objective='reg:squarederror', random_state=42, n_jobs=-1))
])
```

```
best_xgb.fit(X, y) # X and y are the full non-null sets
```

Out [49]:



```
In [50]: # -----
# 12) Impute missing ECPI in the *training* data
# -----
train_missing_idx = df[df['ecpi'].isna()].index
if len(train_missing_idx):
    X_missing_train = df.loc[train_missing_idx].drop(columns=['ecpi', 'churn'])
    df.loc[train_missing_idx, 'ecpi'] = best_xgb.predict(X_missing_train)
    print(f"Filled {len(train_missing_idx)} ECPI values in the training set.")
else:
    print("No missing ECPI rows in training set.")
```

Filled 528 ECPI values in the training set.

```
In [51]: # -----
# Impute missing ECPI in the *test* data – same pattern as the train block
# -----
test_missing_idx = test_df[test_df['ecpi'].isna()].index

if len(test_missing_idx):
    X_missing_test = (
        test_df
        .loc[test_missing_idx]
        .drop(columns=['ecpi', 'churn'], errors='ignore') # keep it identical
    )
    test_df.loc[test_missing_idx, 'ecpi'] = best_xgb.predict(X_missing_test)
    print(f"Filled {len(test_missing_idx)} ECPI values in the test set.")
else:
    print("No missing ECPI rows in test set.")
```

Filled 58 ECPI values in the test set.

```
In [52]: # ----- Quick sanity checks -----
print("\nRemaining nulls → Train:", df['ecpi'].isna().sum(),
      "\t| Test:", test_df['ecpi'].isna().sum())
```

```
print("\nTrain head (after fill):")
display(df.head())

print("\nTest head (after fill):")
display(test_df.head())
```

Remaining nulls → Train: 0 | Test: 0

Train head (after fill):

	country	device_brand	re_install	attribution_event_timestamp	ecpi	lang	curr
0	CZ	Blackview	0	2025-01-06 15:12:35.000	0.279937	CS	
1	AT	samsung	0	2025-01-19 16:50:59.000	0.389500	DE	
2	BG	HONOR	0	2025-01-19 08:46:42.000	0.155800	BU	
3	AZ	samsung	0	2025-01-11 09:10:43.000	0.098094	AZ	
4	AZ	samsung	0	2025-01-16 16:45:16.000	0.083373	AZ	

Test head (after fill):

	test_id	country	device_brand	device_model	re_install	os	attribution_event_t
0	1	SI	samsung	SM-A226B	0	android	2025-01-23 19
1	2	US	samsung	SM-T560NU	0	android	2025-01-07 15
2	3	DZ	samsung	SM-A042F	0	android	2025-01-20 22
3	4	FR	samsung	SM-T500	0	android	2025-01-01 20
4	5	RO	samsung	SM-A145R	0	android	2025-01-04 23

5 rows × 24 columns

MODEL DEVELOPMENT

```
In [53]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, cross_val_score, Strat
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.linear_model import LogisticRegression, LassoCV
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix, classification_report, roc_auc
from xgboost import XGBClassifier
import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.utils.data import TensorDataset, DataLoader
```



```

encoder.fit(train_df[cat_cols])

train_encoded = encoder.transform(train_df[cat_cols])
val_encoded = encoder.transform(val_df[cat_cols])

# Back to DataFrames
train_encoded_df = pd.DataFrame(train_encoded, columns=encoder.get_feature_r
val_encoded_df = pd.DataFrame(val_encoded, columns=encoder.get_feature_names

# Merge back (drop original categorical columns)
train_final = pd.concat([train_df.drop(columns=cat_cols), train_encoded_df],
val_final = pd.concat([val_df.drop(columns=cat_cols), val_encoded_df], axis=

# (Optional) Reset index if you want clean DataFrames
train_final = train_final.reset_index(drop=True)
val_final = val_final.reset_index(drop=True)

# Check result
print(train_final.shape, val_final.shape)
print(train_final.head())

```

```
(44999, 87) (5000, 87)
```

	ecpi	current_gold	totalPowerUp	bonus_cnt	duration	hint1_cnt	\
0	0.489400	920	2	0	16	2	
1	0.153840	1020	3	0	7	0	
2	0.191890	1165	3	0	426	0	
3	0.213389	1115	1	1	1492	1	
4	0.326587	1620	5	9	371	0	

	lvl_no	banner_impr	inter_impr	rewarded_impr	...	campaignid_27	\
0	5	1	0	1	...	0.0	
1	5	2	0	0	...	0.0	
2	5	5	0	4	...	0.0	
3	10	26	0	3	...	1.0	
4	15	32	2	4	...	0.0	

	campaignid_28	campaignid_29	campaignid_30	campaignid_33	campaignid_34	\
0	0.0	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	0.0	0.0	
4	0.0	0.0	0.0	0.0	0.0	

	campaignid_other	partnerid_1	partnerid_4	partnerid_other
0	0.0	1.0	0.0	0.0
1	0.0	0.0	1.0	0.0
2	0.0	0.0	1.0	0.0
3	0.0	0.0	1.0	0.0
4	0.0	0.0	1.0	0.0

```
[5 rows x 87 columns]
```

```

In [57]: from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import PolynomialFeatures
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis

```

```

from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier

from sklearn.metrics import (
    confusion_matrix, classification_report, roc_auc_score, roc_curve, accuracy_score,
    precision_score, recall_score, f1_score
)
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np

# Prepare X, y
X_train = train_final.drop(columns=['churn'])
y_train = train_final['churn']
X_val = val_final.drop(columns=['churn'])
y_val = val_final['churn']

results = {}

def evaluate_model(name, model, X_val, y_val, proba=None):
    y_pred = model.predict(X_val)
    cm = confusion_matrix(y_val, y_pred)
    acc = accuracy_score(y_val, y_pred)
    prec = precision_score(y_val, y_pred)
    rec = recall_score(y_val, y_pred)
    f1 = f1_score(y_val, y_pred)
    # Specificity: TN / (TN + FP)
    tn, fp, fn, tp = cm.ravel()
    spec = tn / (tn + fp)
    if proba is None:
        proba = model.predict_proba(X_val)[:, 1]
    auc = roc_auc_score(y_val, proba)
    fpr, tpr, thresholds = roc_curve(y_val, proba)
    print(f"\n--- {name} ---")
    print(f"Confusion Matrix:\n{cm}")
    print(f"Accuracy: {acc:.4f}")
    print(f"Precision: {prec:.4f}")
    print(f"Recall: {rec:.4f}")
    print(f"Specificity: {spec:.4f}")
    print(f"F1 Score: {f1:.4f}")
    print(f"AUC: {auc:.4f}")
    print(classification_report(y_val, y_pred))
    # Store results
    results[name] = dict(
        accuracy=acc, precision=prec, recall=rec, specificity=spec, f1=f1, auc=auc
    )
    # Plot ROC
    plt.plot(fpr, tpr, label=f"{name} (AUC = {auc:.3f})")

plt.figure(figsize=(10, 8))

```

Out[57]: <Figure size 1000x800 with 0 Axes>

<Figure size 1000x800 with 0 Axes>

```
In [58]: # Logistic Regression
log_reg = LogisticRegression(max_iter=1000)
log_reg.fit(X_train, y_train)
evaluate_model('Logistic Regression', log_reg, X_val, y_val)
```

--- Logistic Regression ---

Confusion Matrix:

```
[[ 416  595]
 [ 145 3844]]
```

Accuracy: 0.8520

Precision: 0.8660

Recall: 0.9637

Specificity: 0.4115

F1 Score: 0.9122

AUC: 0.8136

	precision	recall	f1-score	support
0	0.74	0.41	0.53	1011
1	0.87	0.96	0.91	3989
accuracy			0.85	5000
macro avg	0.80	0.69	0.72	5000
weighted avg	0.84	0.85	0.83	5000

/Users/mustafaercengizmacbooku/miniforge3/lib/python3.12/site-packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):


STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```



```
In [59]: # Polynomial Logistic Regression (degree=2 for illustration)
poly = PolynomialFeatures(degree=2, include_bias=False)
X_train_poly = poly.fit_transform(X_train)
X_val_poly = poly.transform(X_val)
log_reg_poly = LogisticRegression(max_iter=1000)
log_reg_poly.fit(X_train_poly, y_train)
evaluate_model('Poly Logistic Regression', log_reg_poly, X_val_poly, y_val)
```

/Users/mustafaercengizmacbooku/miniforge3/lib/python3.12/site-packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
 Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
 n_iter_i = _check_optimize_result(

--- Poly Logistic Regression ---

Confusion Matrix:

```
[[ 586  425]
```

```
 [ 586 3403]]
```

Accuracy: 0.7978

Precision: 0.8890


Recall: 0.8531

Specificity: 0.5796

F1 Score: 0.8707

AUC: 0.7229

	precision	recall	f1-score	support
0	0.50	0.58	0.54	1011
1	0.89	0.85	0.87	3989
accuracy			0.80	5000
macro avg	0.69	0.72	0.70	5000
weighted avg	0.81	0.80	0.80	5000



```
In [60]: # Linear Discriminant Analysis (LDA)
lda = LinearDiscriminantAnalysis()
lda.fit(X_train, y_train)
evaluate_model('LDA', lda, X_val, y_val)
```

--- LDA ---

Confusion Matrix:

```
[[ 357  654]
```

```
 [  92 3897]]
```

Accuracy: 0.8508

Precision: 0.8563


Recall: 0.9769

Specificity: 0.3531

F1 Score: 0.9126

AUC: 0.8434

	precision	recall	f1-score	support
0	0.80	0.35	0.49	1011
1	0.86	0.98	0.91	3989
accuracy			0.85	5000
macro avg	0.83	0.67	0.70	5000
weighted avg	0.84	0.85	0.83	5000



```
In [61]: # Decision Tree
dt = DecisionTreeClassifier(random_state=42)
dt.fit(X_train, y_train)
evaluate_model('Decision Tree', dt, X_val, y_val)
```

--- Decision Tree ---

Confusion Matrix:

```
[[ 516  495]
 [ 523 3466]]
```

Accuracy: 0.7964

Precision: 0.8750


Recall: 0.8689

Specificity: 0.5104

F1 Score: 0.8719

AUC: 0.6896

	precision	recall	f1-score	support
0	0.50	0.51	0.50	1011
1	0.88	0.87	0.87	3989
accuracy			0.80	5000
macro avg	0.69	0.69	0.69	5000
weighted avg	0.80	0.80	0.80	5000



```
In [62]: # Random Forest
rf = RandomForestClassifier(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)
evaluate_model('Random Forest', rf, X_val, y_val)
```

--- Random Forest ---

Confusion Matrix:

```
[[ 516  495]
 [ 213 3776]]
```

Accuracy: 0.8584

Precision: 0.8841


Recall: 0.9466

Specificity: 0.5104

F1 Score: 0.9143

AUC: 0.8588

	precision	recall	f1-score	support
0	0.71	0.51	0.59	1011
1	0.88	0.95	0.91	3989
accuracy			0.86	5000
macro avg	0.80	0.73	0.75	5000
weighted avg	0.85	0.86	0.85	5000



```
In [63]: # XGBoost
xgb = XGBClassifier(use_label_encoder=False, eval_metric='logloss', random_s
xgb.fit(X_train, y_train)
evaluate_model('XGBoost', xgb, X_val, y_val)
```

```
/Users/mustafaercengizmacbooku/miniforge3/lib/python3.12/site-packages/xgboo
st/core.py:158: UserWarning: [09:55:20] WARNING: /Users/runner/work/xgboost/
xgboost/src/learner.cc:740:
```

```
Parameters: { "use_label_encoder" } are not used.
```


```
warnings.warn(smsg, UserWarning)
```

```

--- XGBoost ---
Confusion Matrix:
[[ 512  499]
 [ 216 3773]]
Accuracy: 0.8570
Precision: 0.8832
Recall: 0.9459
Specificity: 0.5064
F1 Score: 0.9134
AUC: 0.8597

```

	precision	recall	f1-score	support
0	0.70	0.51	0.59	1011
1	0.88	0.95	0.91	3989
accuracy			0.86	5000
macro avg	0.79	0.73	0.75	5000
weighted avg	0.85	0.86	0.85	5000



```

In [64]: # Create results table
results_df = pd.DataFrame(results).T
print("\nSummary Table:")
display(results_df[['accuracy', 'precision', 'recall', 'specificity', 'f1',

```

Summary Table:

	accuracy	precision	recall	specificity	f1	auc
Logistic Regression	0.852	0.865961	0.96365	0.411474	0.912197	0.813624
Poly Logistic Regression	0.7978	0.888976	0.853096	0.579624	0.870666	0.722868
LDA	0.8508	0.856295	0.976937	0.353116	0.912646	0.843357
Decision Tree	0.7964	0.875032	0.868889	0.510386	0.87195	0.689638
Random Forest	0.8584	0.884102	0.946603	0.510386	0.914286	0.858849
XGBoost	0.857	0.883193	0.945851	0.506429	0.913449	0.859699

Testing New Things

```
In [65]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, cross_val_score, Strat
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.linear_model import LogisticRegression, LassoCV
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import (
    confusion_matrix, classification_report, roc_auc_score,
    accuracy_score, precision_score, recall_score, f1_score, roc_curve
)
from xgboost import XGBClassifier
import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.utils.data import TensorDataset, DataLoader
import matplotlib.pyplot as plt

# Copy data
df_model = df.copy()
df_test = test_df.copy()

# Ensure datetime format
df_model['attribution_event_timestamp'] = pd.to_datetime(df_model['attribution_event_timestamp'])

# Extract time-based features
df_model['week'] = df_model['attribution_event_timestamp'].dt.isocalendar().week
df_model['hour'] = df_model['attribution_event_timestamp'].dt.hour
df_model['day'] = df_model['attribution_event_timestamp'].dt.dayofweek # 0

# Preview
df_model.head()
```

```
Out [65]:
```

	country	device_brand	re_install	attribution_event_timestamp	ecpi	lang	cu
0	CZ	Blackview	0	2025-01-06 15:12:35	0.279937	CS	
1	AT	samsung	0	2025-01-19 16:50:59	0.389500	DE	
2	BG	HONOR	0	2025-01-19 08:46:42	0.155800	BU	
3	AZ	samsung	0	2025-01-11 09:10:43	0.098094	AZ	
4	AZ	samsung	0	2025-01-16 16:45:16	0.083373	AZ	

5 rows × 21 columns

```
In [66]: # Step 1: Drop timestamp and lang columns
df_model = df_model.drop(columns=["attribution_event_timestamp", "lang"])

# Step 2: Define columns to one-hot encode
OHE_COLS = ["campaignid", "partnerid", "week", "hour", "day", "re_install", "

# Step 3: Apply one-hot encoding using pandas
df_model = pd.get_dummies(df_model, columns=OHE_COLS, drop_first=True)

# Quick check
print("Final shape after encoding:", df_model.shape)
df_model.head()
```

Final shape after encoding: (49999, 482)

```
Out [66]:
```

	ecpi	current_gold	totalPowerUp	bonus_cnt	duration	hint1_cnt	lvl_no	ban
0	0.279937	120	3	17	3989	37	26	
1	0.389500	65	1	41	8427	46	62	
2	0.155800	1700	5	7	2674	0	36	
3	0.098094	340	3	7	512	28	25	
4	0.083373	660	1	0	978	11	14	

5 rows × 482 columns

```
In [67]: import pandas as pd
import numpy as np
from sklearn.model_selection import StratifiedKFold, train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import (
    accuracy_score, precision_score, recall_score, f1_score,
    roc_auc_score, confusion_matrix, roc_curve
)
from xgboost import XGBClassifier
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler
```

```

from sklearn.pipeline import Pipeline

# Assume df_model is already preprocessed as instructed
# Split features and target
X = df_model.drop(columns=["churn"])
y = df_model["churn"]

# Train/Val split 90/10
X_trainval, X_val, y_trainval, y_val = train_test_split(X, y, test_size=0.1,

# Models
models = {
    "Logistic": LogisticRegression(max_iter=1000),
    "Lasso": LogisticRegression(penalty='l1', solver='liblinear', max_iter=1
    "DecisionTree": DecisionTreeClassifier(),
    "LDA": LinearDiscriminantAnalysis(),
    "RandomForest": RandomForestClassifier(),
    "XGBoost": XGBClassifier(use_label_encoder=False, eval_metric='logloss')
}

# Cross-validation setup
cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
results = {}

```

```

In [69]: for name, model in models.items():
    accs, precs, recs, f1s, rocs = [], [], [], [], []
    tpr_last, fpr_last = None, None

    for train_idx, test_idx in cv.split(X_trainval, y_trainval):
        X_tr, X_ts = X_trainval.iloc[train_idx], X_trainval.iloc[test_idx]
        y_tr, y_ts = y_trainval.iloc[train_idx], y_trainval.iloc[test_idx]

        pipe = Pipeline([
            ("scaler", StandardScaler()),
            ("clf", model)
        ])

        pipe.fit(X_tr, y_tr)
        y_pred = pipe.predict(X_ts)
        y_prob = pipe.predict_proba(X_ts)[:, 1]

        accs.append(accuracy_score(y_ts, y_pred))
        precs.append(precision_score(y_ts, y_pred))
        recs.append(recall_score(y_ts, y_pred))
        f1s.append(f1_score(y_ts, y_pred))
        rocs.append(roc_auc_score(y_ts, y_prob))

        fpr_last, tpr_last, _ = roc_curve(y_ts, y_prob)

    results[name] = {
        "accuracy": np.mean(accs),
        "precision": np.mean(precs),
        "recall": np.mean(recs),
        "f1_score": np.mean(f1s),
        "roc_auc": np.mean(rocs),
        "conf_matrix": confusion_matrix(y_ts, y_pred),

```



```

        "fpr": fpr_last,
        "tpr": tpr_last
    }

# Plot ROC curves
plt.figure(figsize=(10, 8))
for name, res in results.items():
    plt.plot(res["fpr"], res["tpr"], label=f"{name} (AUC = {res['roc_auc']:.2f})")
plt.plot([0, 1], [0, 1], 'k--')
plt.title("ROC Curve Comparison")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend()
plt.grid()
plt.tight_layout()
plt.show()

```

```

/Users/mustafaercengizmacbooku/miniforge3/lib/python3.12/site-packages/xgboost/core.py:158: UserWarning: [10:52:16] WARNING: /Users/runner/work/xgboost/xgboost/src/learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

```

```

    warnings.warn(smsg, UserWarning)
/Users/mustafaercengizmacbooku/miniforge3/lib/python3.12/site-packages/xgboost/core.py:158: UserWarning: [10:52:17] WARNING: /Users/runner/work/xgboost/xgboost/src/learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

```

```

    warnings.warn(smsg, UserWarning)
/Users/mustafaercengizmacbooku/miniforge3/lib/python3.12/site-packages/xgboost/core.py:158: UserWarning: [10:52:18] WARNING: /Users/runner/work/xgboost/xgboost/src/learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

```

```

    warnings.warn(smsg, UserWarning)
/Users/mustafaercengizmacbooku/miniforge3/lib/python3.12/site-packages/xgboost/core.py:158: UserWarning: [10:52:19] WARNING: /Users/runner/work/xgboost/xgboost/src/learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

```

```


    warnings.warn(smsg, UserWarning)
/Users/mustafaercengizmacbooku/miniforge3/lib/python3.12/site-packages/xgboost/core.py:158: UserWarning: [10:52:21] WARNING: /Users/runner/work/xgboost/xgboost/src/learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

```

```

    warnings.warn(smsg, UserWarning)

```



```
In [70]: # Create a summary table excluding FPR/TPR and confusion matrix
summary_df = pd.DataFrame({
    model: {
        "Accuracy": round(metrics["accuracy"], 4),
        "Precision": round(metrics["precision"], 4),
        "Recall": round(metrics["recall"], 4),
        "F1 Score": round(metrics["f1_score"], 4),
        "ROC AUC": round(metrics["roc_auc"], 4)
    }
    for model, metrics in results.items()
}).T

summary_df
```

Out [70]:

	Accuracy	Precision	Recall	F1 Score	ROC AUC
Logistic	0.8543	0.8654	0.9676	0.9137	0.8479
Lasso	0.8544	0.8654	0.9677	0.9137	0.8491
DecisionTree	0.7929	0.8722	0.8669	0.8696	0.6851
LDA	0.8466	0.8521	0.9769	0.9102	0.8460
RandomForest	0.8579	0.8796	0.9518	0.9143	0.8664
XGBoost	0.8539	0.8780	0.9484	0.9118	0.8656

Doing it on Test Set

In [73]: `df_test.head()`

Out [73]:

	test_id	country	device_brand	device_model	re_install	os	attribution_event
0	1	SI	samsung	SM-A226B	0	android	2025-01-23
1	2	US	samsung	SM-T560NU	0	android	2025-01-07
2	3	DZ	samsung	SM-A042F	0	android	2025-01-20
3	4	FR	samsung	SM-T500	0	android	2025-01-01
4	5	RO	samsung	SM-A145R	0	android	2025-01-04

5 rows x 24 columns

```
In [74]: import pandas as pd
from sklearn.ensemble import RandomForestClassifier
from sklearn.preprocessing import OneHotEncoder

# Load data (assuming df and test_df are already loaded)
df_model = df.copy()
test_df_processed = test_df.copy()

# Preprocess training data
df_model['attribution_event_timestamp'] = pd.to_datetime(df_model['attribution_event_timestamp'])
df_model['week'] = df_model['attribution_event_timestamp'].dt.isocalendar().week
df_model['hour'] = df_model['attribution_event_timestamp'].dt.hour
df_model['day'] = df_model['attribution_event_timestamp'].dt.dayofweek
df_model = df_model.drop(columns=['attribution_event_timestamp', 'lang'])

cat_cols = ['country', 'device_brand', 're_install', 'campaignid', 'partnerid']

# Handle rare categories in training data
for col in cat_cols:
    df_model[col] = df_model[col].astype(str)
    freq = df_model[col].value_counts(normalize=True)
```

```
rare_cats = freq[freq < 0.01].index
df_model[col] = df_model[col].replace(rare_cats, 'other')
```

```
In [75]: # Preprocess test data
test_df_processed['attribution_event_timestamp'] = pd.to_datetime(test_df_pr
test_df_processed['week'] = test_df_processed['attribution_event_timestamp']
test_df_processed['hour'] = test_df_processed['attribution_event_timestamp']
test_df_processed['day'] = test_df_processed['attribution_event_timestamp'].
test_df_processed = test_df_processed.drop(columns=['attribution_event_times

# Apply rare categories from training data to test data
for col in cat_cols:
    test_df_processed[col] = test_df_processed[col].astype(str)
    train_cats = df_model[col].unique()
    test_df_processed[col] = test_df_processed[col].apply(lambda x: x if x i

# One-Hot Encoding
encoder = OneHotEncoder(handle_unknown='ignore', sparse_output=False)
encoder.fit(df_model[cat_cols])

X_train_encoded = encoder.transform(df_model[cat_cols])
X_test_encoded = encoder.transform(test_df_processed[cat_cols])
```

```
In [78]: # Create DataFrames
encoded_cols = encoder.get_feature_names_out(cat_cols)
X_train_encoded_df = pd.DataFrame(X_train_encoded, columns=encoded_cols, inc
X_test_encoded_df = pd.DataFrame(X_test_encoded, columns=encoded_cols, index

# Combine with numerical features
numerical_cols = df_model.columns.difference(cat_cols + ['churn']).tolist()
X_train_final = pd.concat([df_model[numerical_cols], X_train_encoded_df], ax
X_test_final = pd.concat([test_df_processed[numerical_cols], X_test_encoded

# Align test columns with training columns
X_test_final = X_test_final.reindex(columns=X_train_final.columns, fill_valu

# Train model on entire training data
y_train = df_model['churn']
best_model = RandomForestClassifier(n_estimators=100, random_state=42)
best_model.fit(X_train_final, y_train)
```

```
Out[78]: ▼ RandomForestClassifier ⓘ ?
RandomForestClassifier(random_state=42)
```

```
In [79]: # Predict and add to test data
test_predictions = best_model.predict(X_test_final)
test_df['churn'] = test_predictions

# Output (if needed)
test_df[['test_id', 'churn']].to_csv('predictions.csv', index=False)
```

```
In [80]: import pandas as pd
import numpy as np
```

```

from xgboost import XGBClassifier
from sklearn.model_selection import cross_val_score, StratifiedKFold
from sklearn.metrics import roc_auc_score, classification_report

# 1) ----- fit XGBoost on the same matrix -----
xgb_model = XGBClassifier(
    n_estimators      = 350,          # ↑ trees → usually better with tabular data
    learning_rate     = 0.05,
    max_depth         = 6,
    subsample         = 0.9,
    colsample_bytree  = 0.8,
    eval_metric       = "auc",
    random_state      = 42,
    n_jobs            = -1
)

xgb_model.fit(X_train_final, y_train)

# 2) ----- predictions & file -----
xgb_preds = xgb_model.predict(X_test_final)

out_xgb = pd.DataFrame({
    "test_id": test_df["test_id"], # adapt if your id column is named differently
    "churn"   : xgb_preds
})
out_xgb.to_csv("prediction_xgb.csv", index=False)
print("✓ Wrote prediction_xgb.csv")

# 3) ----- simple head-to-head comparison on test preds -----
rf_preds = test_predictions # from your RandomForest model
cmp = pd.DataFrame({"rf_pred": rf_preds, "xgb_pred": xgb_preds})
cmp["agree"] = cmp.rf_pred == cmp.xgb_pred

print("\n== Agreement on TEST set ==")
print(cmp["agree"].value_counts(normalize=True).rename({True:"agree", False:"disagree"}))
print("\nCounts predicted as churn (1):")
print({
    "RandomForest": int((cmp.rf_pred==1).sum()),
    "XGBoost"      : int((cmp.xgb_pred==1).sum())
})

# 4) ----- extra: 5-fold CV AUC on TRAIN -----
cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
rf_auc = cross_val_score(best_model, X_train_final, y_train, cv=cv, scoring="roc_auc")
xgb_auc = cross_val_score(xgb_model, X_train_final, y_train, cv=cv, scoring="roc_auc")

print(f"\n=== 5-fold CV AUC (train) ===")
print(f"RandomForest : {rf_auc.mean():.4f} ± {rf_auc.std():.4f}")
print(f"XGBoost       : {xgb_auc.mean():.4f} ± {xgb_auc.std():.4f}")

# 5) ----- optional: full classification report on TRAIN -----
print("\n== Full classification report on training set (fit-on-full) ==")
for name, clf in [("RandomForest", best_model), ("XGBoost", xgb_model)]:
    print(f"\n{name}")
    y_hat = clf.predict(X_train_final)
    print(classification_report(y_train, y_hat, digits=3))

```

✓ Wrote prediction_xgb.csv

== Agreement on TEST set ==

agree

agree 96.75%

disagree 3.25%

Name: proportion, dtype: object

Counts predicted as churn (1):

{'RandomForest': 5110, 'XGBoost': 5129}

=== 5-fold CV AUC (train) ===

RandomForest : 0.8685 ± 0.0014

XGBoost : 0.8724 ± 0.0008

== Full classification report on training set (fit-on-full) ==

RandomForest

	precision	recall	f1-score	support
0	1.000	1.000	1.000	10179
1	1.000	1.000	1.000	39820
accuracy			1.000	49999
macro avg	1.000	1.000	1.000	49999
weighted avg	1.000	1.000	1.000	49999

XGBoost

	precision	recall	f1-score	support
0	0.822	0.553	0.661	10179
1	0.895	0.969	0.930	39820
accuracy			0.885	49999
macro avg	0.858	0.761	0.796	49999
weighted avg	0.880	0.885	0.876	49999

INFERENCES

```
In [81]: # -----  
# 1. Prep  
# -----  
import pandas as pd  
import matplotlib.pyplot as plt  
from sklearn.linear_model import LogisticRegression  
from sklearn.preprocessing import StandardScaler  
from sklearn.pipeline import make_pipeline  
  
TOP_N = 25                   # how many features to display  
FIG_W, FIG_H = 6, 7        # figure size  
  
# -----
```

```

# 2. Random-Forest feature importance
# -----
rf_imp = (pd.Series(best_model.feature_importances_,
                    index=X_train_final.columns)
          .sort_values(ascending=False))

print("\n=== Random-Forest - Top {} features ===".format(TOP_N))
display(rf_imp.head(TOP_N).to_frame("Importance"))

plt.figure(figsize=(FIG_W, FIG_H))
rf_imp.head(TOP_N).iloc[::-1].plot(kind="barh")
plt.title("Random Forest - Top {} feature importances".format(TOP_N))
plt.xlabel("Gini importance")
plt.tight_layout()
plt.show()

# -----
# 3. XGBoost feature importance
# -----
xgb_imp = (pd.Series(xgb_model.feature_importances_,
                    index=X_train_final.columns)
          .sort_values(ascending=False))

print("\n=== XGBoost - Top {} features ===".format(TOP_N))
display(xgb_imp.head(TOP_N).to_frame("Gain"))

plt.figure(figsize=(FIG_W, FIG_H))
xgb_imp.head(TOP_N).iloc[::-1].plot(kind="barh")
plt.title("XGBoost - Top {} feature importances".format(TOP_N))
plt.xlabel("Split-gain importance")
plt.tight_layout()
plt.show()

# -----
# 4. Logistic-Regression coefficients (signed!)
# -----
log_reg_full = make_pipeline(
    StandardScaler(with_mean=False),      # sparse matrix friendly
    LogisticRegression(max_iter=2000, n_jobs=-1)
).fit(X_train_final, y_train)

coef = pd.Series(log_reg_full[-1].coef_[0], index=X_train_final.columns)
coef_table = (pd.DataFrame({
    "coef"      : coef,
    "abs_coef"  : coef.abs()
}))
    .sort_values("abs_coef", ascending=False)
    .head(TOP_N))

print("\n=== Logistic Regression - Top {} coefficients ===".format(TOP_N))
display(coef_table)

plt.figure(figsize=(FIG_W, FIG_H))
coef_table.sort_values("coef").plot(

```


```

    y="coef", kind="barh", legend=False
)
plt.title("Logistic Regression – Top signed coefficients")
plt.xlabel("Coefficient (positive → ↑ churn prob)")
plt.tight_layout()
plt.show()

```


=== Random-Forest – Top 25 features ===

	Importance
duration	0.120842
lvl_no	0.103946
inter_impr	0.101383
banner_impr	0.097367
bonus_cnt	0.083007
current_gold	0.054182
rewarded_impr	0.048726
hint1_cnt	0.046803
ecpi	0.043514
totalPowerUp	0.021075
device_brand_samsung	0.008762
day_3	0.006607
device_brand_Xiaomi	0.006546
day_6	0.006261
day_4	0.006230
day_2	0.006186
week_2	0.006136
day_5	0.005947
week_4	0.005934
day_0	0.005610
week_1	0.005582
week_5	0.005263
country_FR	0.005084
campaignid_25	0.005076
hour_21	0.004976



=== XGBoost - Top 25 features ===

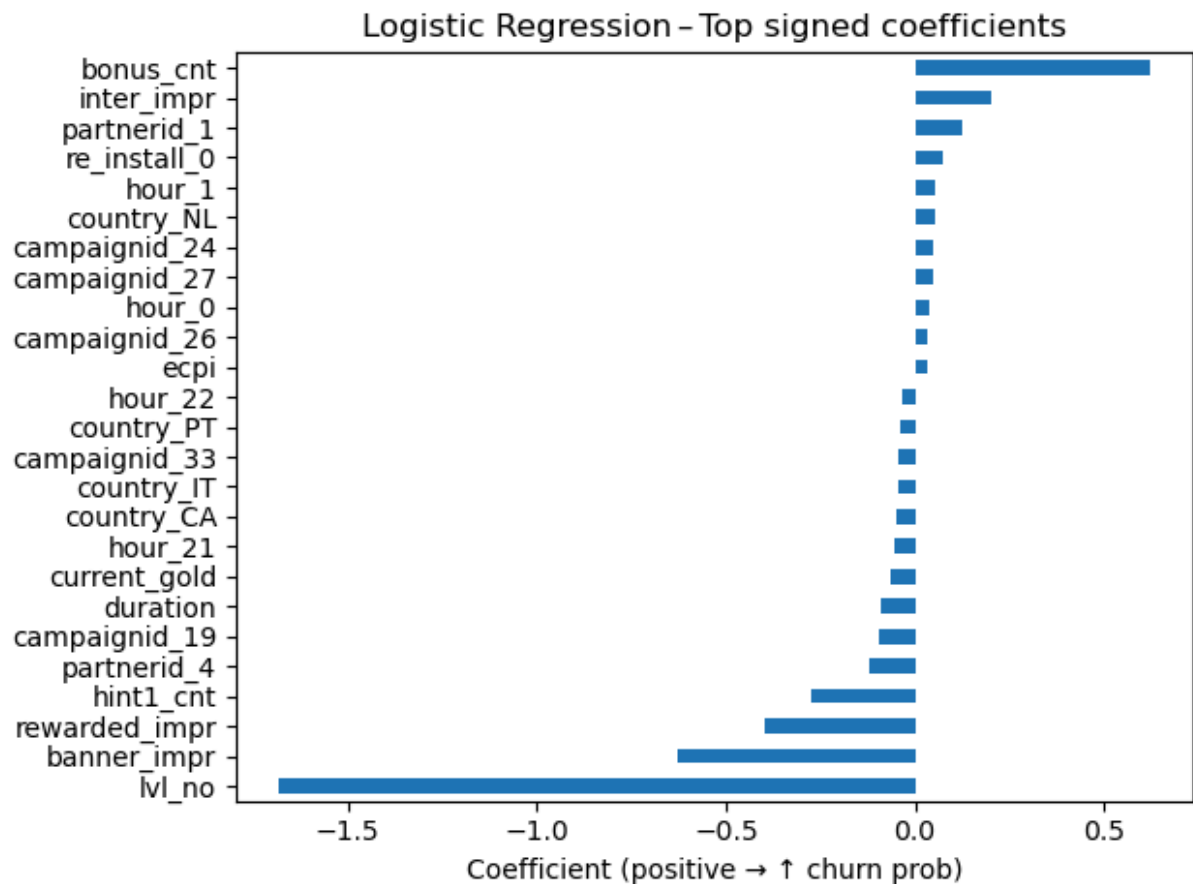
	Gain
inter_impr	0.179147
duration	0.066264
lvl_no	0.032028
partnerid_4	0.013612
partnerid_1	0.011331
hour_21	0.010840
banner_impr	0.010390
campaignid_23	0.010113
rewarded_impr	0.010081
campaignid_33	0.009483
campaignid_10	0.009131
country_NL	0.009113
hour_1	0.009009
campaignid_15	0.008897
country_ID	0.008752
week_3	0.008680
campaignid_29	0.008545
country_SK	0.008542
country_CA	0.008507
campaignid_19	0.008435
hour_6	0.008421
country_BE	0.008402
campaignid_28	0.008324
country_RO	0.008320
device_brand_motorola	0.008298



=== Logistic Regression - Top 25 coefficients ===

	coef	abs_coef
lvl_no	-1.679933	1.679933
banner_impr	-0.626174	0.626174
bonus_cnt	0.619268	0.619268
rewarded_impr	-0.395857	0.395857
hint1_cnt	-0.274279	0.274279
inter_impr	0.202118	0.202118
partnerid_1	0.126378	0.126378
partnerid_4	-0.122078	0.122078
campaignid_19	-0.093440	0.093440
duration	-0.090270	0.090270
re_install_0	0.074052	0.074052
current_gold	-0.063737	0.063737
hour_21	-0.052000	0.052000
hour_1	0.051898	0.051898
country_NL	0.051585	0.051585
campaignid_24	0.050093	0.050093
country_CA	-0.047329	0.047329
country_IT	-0.046108	0.046108
campaignid_27	0.045852	0.045852
campaignid_33	-0.045255	0.045255
country_PT	-0.039227	0.039227
hour_0	0.035568	0.035568
campaignid_26	0.035061	0.035061
hour_22	-0.034203	0.034203
ecpi	0.033627	0.033627

<Figure size 600x700 with 0 Axes>



NEW FEATURE

```
In [83]: # -----
# 1. Revised helper
# -----
def preprocess_data(df, *, is_train=True, encoder=None):
    """
    Returns
    -----
    X_df      : pd.DataFrame – model-ready features
    encoder   : fitted OneHotEncoder (returned only when is_train=True)
    """
    df = df.copy()

    # ----- time features -----
    df['attribution_event_timestamp'] = pd.to_datetime(df['attribution_event_timestamp'])
    df['week'] = df['attribution_event_timestamp'].dt.isocalendar().week
    df['hour'] = df['attribution_event_timestamp'].dt.hour
    df['day'] = df['attribution_event_timestamp'].dt.dayofweek
    df = df.drop(columns=['attribution_event_timestamp', 'lang'])

    # ----- new engineered features -----
    df['lvl_duration_ratio'] = df['lvl_no'] / (df['duration'] + 1)
    df['log_lvl_duration_ratio'] = np.log1p(df['lvl_duration_ratio'])

    # ----- categorical handling -----
```

```

cat_cols = ['country', 'device_brand', 're_install',
            'campaignid', 'partnerid', 'week', 'hour', 'day']

if is_train:
    # group rares → 'other', THEN fit encoder
    for col in cat_cols:
        df[col] = df[col].astype(str)
        rares = (df[col].value_counts(normalize=True)
                 .loc[lambda s: s < 0.01].index)
        df[col] = df[col].replace(rares, 'other')

    encoder = OneHotEncoder(handle_unknown='ignore',
                             sparse_output=False)
    encoder.fit(df[cat_cols])

else:
    # map unseen cats → 'other' using encoder.categories_
    for idx, col in enumerate(cat_cols):
        df[col] = df[col].astype(str)
        known = set(encoder.categories_[idx])
        df[col] = df[col].where(df[col].isin(known), 'other')

# ----- one-hot encode -----
ohe = encoder.transform(df[cat_cols])
ohe_cols = encoder.get_feature_names_out(cat_cols)
ohe_df = pd.DataFrame(ohe, columns=ohe_cols, index=df.index)

# ----- numeric columns -----
num_cols = ['ecpi', 'current_gold', 'totalPowerUp', 'bonus_cnt',
            'duration', 'hint1_cnt', 'lvl_no', 'banner_impr',
            'inter_impr', 'rewarded_impr',
            'lvl_duration_ratio', 'log_lvl_duration_ratio']

X_df = pd.concat([df[num_cols].reset_index(drop=True),
                  ohe_df.reset_index(drop=True)], axis=1)

return (X_df, encoder) if is_train else (X_df, None)

# -----
# 2. Build train + test matrices
# -----
X_train_full, encoder = preprocess_data(df, is_train=True)
y_train_full = df['churn']

X_test_proc, _ = preprocess_data(test_df, is_train=False, encoder=enc

# guarantee identical column order (in case the fit order ever changes)
X_test_proc = X_test_proc.reindex(columns=X_train_full.columns, fill_value=0)

# -----
# 3. Fit the tuned Random-Forest
# -----
best_model = RandomForestClassifier(
    n_estimators = 200,

```

```

    max_depth      = 12,
    min_samples_leaf = 5,
    class_weight    = 'balanced',
    random_state    = 42
).fit(X_train_full, y_train_full)

# -----
# 4. Feature importance + prediction
# -----

importances = (pd.Series(best_model.feature_importances_,
                        index=X_train_full.columns)
               .sort_values(ascending=False))

print("\nTop 10 features:\n", importances.head(10))

test_df['churn'] = best_model.predict(X_test_proc)
test_df[['test_id', 'churn']].to_csv('predictions_with_new_feats.csv', index=False)
print("✓ New predictions saved to predictions_with_new_feats.csv")

```

Top 10 features:

inter_impr	0.148975
lvl_no	0.148120
duration	0.132121
banner_impr	0.112467
lvl_duration_ratio	0.086510
log_lvl_duration_ratio	0.080682
bonus_cnt	0.073746
hint1_cnt	0.046696
rewarded_impr	0.044342
current_gold	0.034803
dtype:	float64

✓ New predictions saved to predictions_with_new_feats.csv

```

In [84]: # -----
# Cross-validated performance of the new feature set
# -----

import numpy as np
import pandas as pd
from sklearn.model_selection import StratifiedKFold
from sklearn.metrics import (accuracy_score, precision_score, recall_score,
                             f1_score, roc_auc_score, confusion_matrix)

cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

fold_metrics = {
    "fold": [], "accuracy": [], "precision": [], "recall": [],
    "specificity": [], "f1": [], "roc_auc": []
}

for fold, (idx_tr, idx_val) in enumerate(cv.split(X_train_full, y_train_full)):
    X_tr, X_val = X_train_full.iloc[idx_tr], X_train_full.iloc[idx_val]
    y_tr, y_val = y_train_full.iloc[idx_tr], y_train_full.iloc[idx_val]

    # fresh clone each fold
    clf = RandomForestClassifier(

```

```

        n_estimators      = 200,
        max_depth         = 12,
        min_samples_leaf  = 5,
        class_weight       = 'balanced',
        random_state      = 42
    ).fit(X_tr, y_tr)

    y_pred = clf.predict(X_val)
    y_prob = clf.predict_proba(X_val)[:, 1]

    cm = confusion_matrix(y_val, y_pred)
    tn, fp, fn, tp = cm.ravel()

    fold_metrics["fold"].append(fold)
    fold_metrics["accuracy"].append(accuracy_score(y_val, y_pred))
    fold_metrics["precision"].append(precision_score(y_val, y_pred))
    fold_metrics["recall"].append(recall_score(y_val, y_pred))
    fold_metrics["specificity"].append(tn / (tn + fp))
    fold_metrics["f1"].append(f1_score(y_val, y_pred))
    fold_metrics["roc_auc"].append(roc_auc_score(y_val, y_prob))

# -----
# Pretty summary
# -----
cv_summary = pd.DataFrame(fold_metrics)
means = cv_summary.mean(numeric_only=True).rename("mean").to_frame().T
print("\n=== 5-Fold CV - per-fold metrics ===")
display(cv_summary.round(4))

print("\n=== 5-Fold CV - averaged ===")
display(means.round(4))

```

=== 5-Fold CV - per-fold metrics ===

	fold	accuracy	precision	recall	specificity	f1	roc_auc
0	1	0.8102	0.9268	0.8270	0.7446	0.8741	0.8714
1	2	0.8102	0.9256	0.8282	0.7397	0.8742	0.8712
2	3	0.8101	0.9245	0.8292	0.7353	0.8743	0.8704
3	4	0.8136	0.9254	0.8331	0.7372	0.8768	0.8753
4	5	0.8120	0.9233	0.8331	0.7292	0.8759	0.8734

=== 5-Fold CV - averaged ===

	fold	accuracy	precision	recall	specificity	f1	roc_auc
mean	3.0	0.8112	0.9251	0.8301	0.7372	0.8751	0.8723

In []: